

An introduction to deep-learning-based methods for optimization and control of PDEs

Francisco Periago

Universidad Politécnica de Cartagena

https://github.com/fperiago/pinn_deeponet_for_beginners

**XXI Jacques-Louis Lions Hispano-French School on Numerical
Simulation in Physics and Engineering**

Ciudad Real, July, 7-11, 2025



Universidad
Politécnica
de Cartagena



EUROPEAN
UNIVERSITY OF
TECHNOLOGY



- **Part I: Machine Learning basis.** Where and why Machine-Learning-based methods may be useful in the numerical approximation of PDEs-based models?

- **Part I: Machine Learning basis.** Where and why Machine-Learning-based methods may be useful in the numerical approximation of PDEs-based models?
- **Part II: Physics Informed Neural Networks (PINNs)** A machine-learning-based method for solving direct, inverse, control, etc., PDEs-based problems for fixed initial and/or boundary conditions

- **Part I: Machine Learning basis.** Where and why Machine-Learning-based methods may be useful in the numerical approximation of PDEs-based models?
- **Part II: Physics Informed Neural Networks (PINNs)** A machine-learning-based method for solving direct, inverse, control, etc., PDEs-based problems for fixed initial and/or boundary conditions
- **Part III: Deep Operator Network (DeepONet).** Learning the solution map of the above problems

- **Part I: Machine Learning basis.** Where and why Machine-Learning-based methods may be useful in the numerical approximation of PDEs-based models?
- **Part II: Physics Informed Neural Networks (PINNs)** A machine-learning-based method for solving direct, inverse, control, etc., PDEs-based problems for fixed initial and/or boundary conditions
- **Part III: Deep Operator Network (DeepONet).** Learning the solution map of the above problems
- **Part IV: Numerical implementation via DeepXDE.** A Python library for scientific machine learning and physics-informed learning

Part I

Machine Learning Basis

Introduction: the set up of supervised learning



Introduction: the set up of supervised learning



Main Goal:

Introduction: the set up of supervised learning



Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, y_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Introduction: the set up of supervised learning



Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, y_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

Introduction: the set up of supervised learning



Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, y_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and

Introduction: the set up of supervised learning



Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, y_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and
- 2 **classification:** f^* takes discrete values.

Introduction: the set up of supervised learning



Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, y_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and
- 2 **classification:** f^* takes discrete values.

Standard procedure for supervised learning(regression)

Introduction: the set up of supervised learning



Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, y_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and
- 2 **classification:** f^* takes discrete values.

Standard procedure for supervised learning(regression)

- 1 Choose a hypothesis space $\mathcal{H}_m$.

Introduction: the set up of supervised learning



Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, y_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and
- 2 **classification:** f^* takes discrete values.

Standard procedure for supervised learning(regression)

- 1 Choose a hypothesis space $\mathcal{H}_m$. *Artificial neural networks* is the model of choice in Machine Learning.

Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, \mathbf{y}_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and
- 2 **classification:** f^* takes discrete values.

Standard procedure for supervised learning(regression)

- 1 Choose a hypothesis space $\mathcal{H}_m$. *Artificial neural networks* is the model of choice in Machine Learning.
- 2 Choose a loss function. If we are interested in fitting the data, a popular choice is the so-called **training error**

$$\hat{\mathcal{R}}_n(f) = \frac{1}{n} \sum_{i=1}^n (f(\theta; \mathbf{x}_i) - f^*(\mathbf{x}_i))^2, \quad f \in \mathcal{H}_m. \quad (1)$$

Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, \mathbf{y}_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and
- 2 **classification:** f^* takes discrete values.

Standard procedure for supervised learning(regression)

- 1 Choose a hypothesis space $\mathcal{H}_m$. *Artificial neural networks* is the model of choice in Machine Learning.
- 2 Choose a loss function. If we are interested in fitting the data, a popular choice is the so-called **training error**

$$\hat{\mathcal{R}}_n(f) = \frac{1}{n} \sum_{i=1}^n (f(\boldsymbol{\theta}; \mathbf{x}_i) - f^*(\mathbf{x}_i))^2, \quad f \in \mathcal{H}_m. \quad (1)$$

- 3 Choose an optimization algorithm for computing the optimal parameters $\boldsymbol{\theta}$ that minimize the loss function.

Main Goal: approximate (as accurately as we can) an unknown function $f^* : \mathbb{R}^d \rightarrow \mathbb{R}^N$ from a dataset $S = \{(\mathbf{x}_i, \mathbf{y}_i = f^*(\mathbf{x}_i)), 1 \leq i \leq n\}$

Two cases:

- 1 **regression:** f^* takes continuous values, and
- 2 **classification:** f^* takes discrete values.

Standard procedure for supervised learning(regression)

- 1 Choose a hypothesis space $\mathcal{H}_m$. *Artificial neural networks* is the model of choice in Machine Learning.
- 2 Choose a loss function. If we are interested in fitting the data, a popular choice is the so-called **training error**

$$\hat{\mathcal{R}}_n(f) = \frac{1}{n} \sum_{i=1}^n (f(\boldsymbol{\theta}; \mathbf{x}_i) - f^*(\mathbf{x}_i))^2, \quad f \in \mathcal{H}_m. \quad (1)$$

- 3 Choose an optimization algorithm for computing the optimal parameters $\boldsymbol{\theta}$ that minimize the loss function.

The overall objective is to minimize the **generalization error**

$$\mathcal{R}(f) = \mathbb{E}_{\mathbf{x} \sim \mathbb{P}} (f(\boldsymbol{\theta}; \mathbf{x}) - f^*(\mathbf{x}))^2, \quad f \in \mathcal{H}_m, \quad (2)$$

with $\mathbb{P}$ the (unknown) distribution of $\mathbf{x}$.

What separates ML from optimization is that we want the **generalization error**, also called the **test error**, to be low as well. The factors determining how well a machine learning algorithm will perform are its ability to:

- Make the training error small, and
- Make the gap between training and test error (**generalization gap**) small.

What separates ML from optimization is that we want the **generalization error**, also called the **test error**, to be low as well. The factors determining how well a machine learning algorithm will perform are its ability to:

- Make the training error small, and
- Make the gap between training and test error (**generalization gap**) small.

These two factors correspond to the so-called:

- **Underfitting**, when the model is not able to obtain a sufficiently low error value on the training set
- **Overfitting**, when the generalization gap is too large.

What separates ML from optimization is that we want the **generalization error**, also called the **test error**, to be low as well. The factors determining how well a machine learning algorithm will perform are its ability to:

- Make the training error small, and
- Make the gap between training and test error (**generalization gap**) small.

These two factors correspond to the so-called:

- **Underfitting**, when the model is not able to obtain a sufficiently low error value on the training set
- **Overfitting**, when the generalization gap is too large.

Capacity is the ability of a ML algorithm to fit a wide variety of functions.

Analyzing the performance of a ML algorithm


What separates ML from optimization is that we want the **generalization error**, also called the **test error**, to be low as well. The factors determining how well a machine learning algorithm will perform are its ability to:

- Make the training error small, and
- Make the gap between training and test error (**generalization gap**) small.

These two factors correspond to the so-called:

- **Underfitting**, when the model is not able to obtain a sufficiently low error value on the training set
- **Overfitting**, when the generalization gap is too large.


Capacity is the ability of a ML algorithm to fit a wide variety of functions.



Hypothesis space: an example

A canonical example of an hypothesis space $\mathcal{H}_m$ (or a neural network architecture) is the so-called **Multi-Layer Perceptron (MLP)**.


Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer



Hypothesis space: an example

A canonical example of an hypothesis space $\mathcal{H}_m$ (or a neural network architecture) is the so-called **Multi-Layer Perceptron (MLP)**.

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer




To each input $\mathbf{x} \in \mathbb{R}^d$ it associates the output $\mathbf{y} = f_m(\mathbf{x}) := \mathbf{x}^m$ defined by

$$\left\{ \begin{array}{ll} \text{input layer:} & \mathbf{x}^0 = \mathbf{x}, \\ \text{hidden layers:} & \mathbf{x}^{k+1} = \sigma(\omega^{k+1} \mathbf{x}^k + b^{k+1}) \quad \text{for } k = 0, 1, \dots, m-2 \\ \text{output layer:} & \mathbf{x}^m = \omega^m \mathbf{x}^{m-1} + b^m, \end{array} \right.$$

Hypothesis space: an example

A canonical example of an hypothesis space $\mathcal{H}_m$ (or a neural network architecture) is the so-called **Multi-Layer Perceptron (MLP)**.

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer



To each input $\mathbf{x} \in \mathbb{R}^d$ it associates the output $\mathbf{y} = f_m(\mathbf{x}) := \mathbf{x}^m$ defined by


$$\begin{cases} \text{input layer:} & \mathbf{x}^0 = \mathbf{x}, \\ \text{hidden layers:} & \mathbf{x}^{k+1} = \sigma(\omega^{k+1}\mathbf{x}^k + \mathbf{b}^{k+1}) \quad \text{for } k = 0, 1, \dots, m-2 \\ \text{output layer:} & \mathbf{x}^m = \omega^m\mathbf{x}^{m-1} + \mathbf{b}^m, \end{cases}$$

■ optimizable parameters θ : **weights** $\omega^k \in \mathbb{R}^{d_{k+1} \times d_k}$ and **biases** $\mathbf{b}^k \in \mathbb{R}^{d_k}$

Hypothesis space: an example

A canonical example of an hypothesis space $\mathcal{H}_m$ (or a neural network architecture) is the so-called **Multi-Layer Perceptron (MLP)**.

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer



To each input $\mathbf{x} \in \mathbb{R}^d$ it associates the output $\mathbf{y} = f_m(\mathbf{x}) := \mathbf{x}^m$ defined by


$$\begin{cases} \text{input layer:} & \mathbf{x}^0 = \mathbf{x}, \\ \text{hidden layers:} & \mathbf{x}^{k+1} = \sigma(\omega^{k+1}\mathbf{x}^k + \mathbf{b}^{k+1}) \quad \text{for } k = 0, 1, \dots, m-2 \\ \text{output layer:} & \mathbf{x}^m = \omega^m\mathbf{x}^{m-1} + \mathbf{b}^m, \end{cases}$$

- optimizable parameters θ : **weights** $\omega^k \in \mathbb{R}^{d_{k+1} \times d_k}$ and **biases** $\mathbf{b}^k \in \mathbb{R}^{d_k}$
- m is the **depth** of the neural network,

Hypothesis space: an example

A canonical example of an hypothesis space $\mathcal{H}_m$ (or a neural network architecture) is the so-called **Multi-Layer Perceptron (MLP)**.

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer



To each input $\mathbf{x} \in \mathbb{R}^d$ it associates the output $\mathbf{y} = f_m(\mathbf{x}) := \mathbf{x}^m$ defined by


$$\begin{cases} \text{input layer:} & \mathbf{x}^0 = \mathbf{x}, \\ \text{hidden layers:} & \mathbf{x}^{k+1} = \sigma(\omega^{k+1}\mathbf{x}^k + b^{k+1}) \quad \text{for } k = 0, 1, \dots, m-2 \\ \text{output layer:} & \mathbf{x}^m = \omega^m\mathbf{x}^{m-1} + b^m, \end{cases}$$

- optimizable parameters θ : **weights** $\omega^k \in \mathbb{R}^{d_{k+1} \times d_k}$ and **biases** $b^k \in \mathbb{R}^{d_k}$
- m is the **depth** of the neural network,
- for any k , the vector $\mathbf{x}^k \in \mathbb{R}^{d_k}$ and d_k is the **width** of the layer k ,

Hypothesis space: an example

A canonical example of an hypothesis space $\mathcal{H}_m$ (or a neural network architecture) is the so-called **Multi-Layer Perceptron (MLP)**.

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer



To each input $\mathbf{x} \in \mathbb{R}^d$ it associates the output $\mathbf{y} = f_m(\mathbf{x}) := \mathbf{x}^m$ defined by

$$\begin{cases} \text{input layer:} & \mathbf{x}^0 = \mathbf{x}, \\ \text{hidden layers:} & \mathbf{x}^{k+1} = \sigma(\omega^{k+1}\mathbf{x}^k + b^{k+1}) \quad \text{for } k = 0, 1, \dots, m-2 \\ \text{output layer:} & \mathbf{x}^m = \omega^m\mathbf{x}^{m-1} + b^m, \end{cases}$$

- optimizable parameters θ : **weights** $\omega^k \in \mathbb{R}^{d_{k+1} \times d_k}$ and **biases** $b^k \in \mathbb{R}^{d_k}$
- m is the **depth** of the neural network,
- for any k , the vector $\mathbf{x}^k \in \mathbb{R}^{d_k}$ and d_k is the **width** of the layer k ,
- σ is a fixed nonlinear **activation function**

Hypothesis space: an example



More on the **activation function**:

Hypothesis space: an example



More on the **activation function**:

By abuse of notation, $\sigma : \mathbb{R}^d \rightarrow \mathbb{R}^d$ is defined component-wise by

$$\sigma(\mathbf{x})_j := \sigma(x_j), \quad 1 \leq j \leq d.$$

Hypothesis space: an example

More on the **activation function**:

By abuse of notation, $\sigma : \mathbb{R}^d \rightarrow \mathbb{R}^d$ is defined component-wise by

$$\sigma(\mathbf{x})_j := \sigma(x_j), \quad 1 \leq j \leq d.$$

Common choices include *rectifiers* such as ReLU: $\sigma(s) = \max\{s, 0\}$, and *sigmoids* such as $\sigma(s) = \tanh(s)$.




Figure: Linear Unit Rectifier (left) and hyperbolic tangent (right).

Algorithm 1: Adaptive with Moment (Adam)

- 1 **Requieres:** learning rate ϵ (0.001) and decreasing rates for the moments:
 ρ_1 y ρ_2 ($\rho_1 = 0.9$, $\rho_2 = 0.999$).

Algorithm 2: Adaptive with Moment (Adam)

- 1 **Requieres:** learning rate ϵ (0.001) and decreasing rates for the moments:
 ρ_1 y ρ_2 ($\rho_1 = 0.9$, $\rho_2 = 0.999$).
- 2 **Initialization:** $\theta^{(0)} = v^{(0)} = w^{(0)} = (0, \dots, 0) \in \mathbb{R}^r$, $k = 0$.

Algorithm 3: Adaptive with Moment (Adam)

- 1 **Requieres:** learning rate ϵ (0.001) and decreasing rates for the moments:
 ρ_1 y ρ_2 ($\rho_1 = 0.9$, $\rho_2 = 0.999$).
- 2 **Initialization:** $\theta^{(0)} = v^{(0)} = w^{(0)} = (0, \dots, 0) \in \mathbb{R}^r$, $k = 0$.
- 3 **while** *stopping condition is not satisfied* **do**
- 4 Select a batch of size s from the training dataset:
 $\{(x^{(k)}_1; y^{(k)}_1), \dots, (x^{(k)}_s; y^{(k)}_s)\}$

Algorithm 4: Adaptive with Moment (Adam)

- 1 **Requieres:** learning rate ϵ (0.001) and decreasing rates for the moments:
 ρ_1 y ρ_2 ($\rho_1 = 0.9$, $\rho_2 = 0.999$).
- 2 **Initialization:** $\theta^{(0)} = v^{(0)} = w^{(0)} = (0, \dots, 0) \in \mathbb{R}^r$, $k = 0$.
- 3 **while** *stopping condition is not satisfied* **do**
- 4 Select a batch of size s from the training dataset:
 $\{(x^{(k)}_1; y^{(k)}_1), \dots, (x^{(k)}_s; y^{(k)}_s)\}$
- 5 Compute an estimation $g^{(k)}$ of the gradient of $\text{loss}(\theta)$:
 $g^{(k)} \leftarrow \frac{1}{s} \sum_{j=1}^s \nabla_{\theta} \text{loss}(\theta; (x^{(k)}_j, y^{(k)}_j))$

Algorithm 5: Adaptive with Moment (Adam)

- 1 **Requieres:** learning rate ϵ (0.001) and decreasing rates for the moments:
 ρ_1 y ρ_2 ($\rho_1 = 0.9$, $\rho_2 = 0.999$).
- 2 **Initialization:** $\theta^{(0)} = v^{(0)} = w^{(0)} = (0, \dots, 0) \in \mathbb{R}^r$, $k = 0$.
- 3 **while** *stopping condition is not satisfied* **do**
- 4 Select a batch of size s from the training dataset:
 $\{(x^{(k)}_1; y^{(k)}_1), \dots, (x^{(k)}_s; y^{(k)}_s)\}$
- 5 Compute an estimation $g^{(k)}$ of the gradient of $\text{loss}(\theta)$:
 $g^{(k)} \leftarrow \frac{1}{s} \sum_{j=1}^s \nabla_{\theta} \text{loss}(\theta; (x^{(k)}_j, y^{(k)}_j))$
- 6 Update biased moments:
 $v^{(k+1)} \leftarrow \rho_1 v^{(k)} + (1 - \rho_1) g^{(k)}$
 $w^{(k+1)} \leftarrow \rho_2 w^{(k)} + (1 - \rho_2) g^{(k)} \odot g^{(k)}$

Algorithm 6: Adaptive with Moment (Adam)

1 **Requieres:** learning rate ϵ (0.001) and decreasing rates for the moments:

ρ_1 y ρ_2 ($\rho_1 = 0.9$, $\rho_2 = 0.999$).

2 **Initialization:** $\theta^{(0)} = v^{(0)} = w^{(0)} = (0, \dots, 0) \in \mathbb{R}^r$, $k = 0$.

3 **while** *stopping condition is not satisfied* **do**

4 Select a batch of size s from the training dataset:

$$\{(x^{(k)_1}; y^{(k)_1}), \dots, (x^{(k)_s}; y^{(k)_s})\}$$

5 Compute an estimation $g^{(k)}$ of the gradient of $\text{loss}(\theta)$:

$$g^{(k)} \leftarrow \frac{1}{s} \sum_{j=1}^s \nabla_{\theta} \text{loss}(\theta; (x^{(k)_j}, y^{(k)_j}))$$

6 Update biased moments:

$$v^{(k+1)} \leftarrow \rho_1 v^{(k)} + (1 - \rho_1) g^{(k)}$$

$$w^{(k+1)} \leftarrow \rho_2 w^{(k)} + (1 - \rho_2) g^{(k)} \odot g^{(k)}$$

7 Correct the biases:

$$\hat{v}^{(k+1)} = \frac{v^{(k+1)}}{1 - \rho_1^{k+1}}, \quad \hat{w}^{(k+1)} = \frac{w^{(k+1)}}{1 - \rho_2^{k+1}}$$

Algorithm 7: Adaptive with Moment (Adam)

1 **Requieres:** learning rate ϵ (0.001) and decreasing rates for the moments:
 ρ_1 y ρ_2 ($\rho_1 = 0.9$, $\rho_2 = 0.999$).

2 **Initialization:** $\theta^{(0)} = v^{(0)} = w^{(0)} = (0, \dots, 0) \in \mathbb{R}^r$, $k = 0$.

3 **while** *stopping condition is not satisfied* **do**

4 Select a batch of size s from the training dataset:

$$\{(x^{(k)_1}; y^{(k)_1}), \dots, (x^{(k)_s}; y^{(k)_s})\}$$

5 Compute an estimation $g^{(k)}$ of the gradient of $\text{loss}(\theta)$:

$$g^{(k)} \leftarrow \frac{1}{s} \sum_{j=1}^s \nabla_{\theta} \text{loss}(\theta; (x^{(k)_j}, y^{(k)_j}))$$

6 Update biased moments:

$$v^{(k+1)} \leftarrow \rho_1 v^{(k)} + (1 - \rho_1) g^{(k)}$$

$$w^{(k+1)} \leftarrow \rho_2 w^{(k)} + (1 - \rho_2) g^{(k)} \odot g^{(k)}$$

7 Correct the biases:

$$\hat{v}^{(k+1)} = \frac{v^{(k+1)}}{1 - \rho_1^{k+1}}, \quad \hat{w}^{(k+1)} = \frac{w^{(k+1)}}{1 - \rho_2^{k+1}}$$


8 Update the optimization variable:

$$\theta_i^{(k+1)} \leftarrow \theta_i^{(k)} - \epsilon \frac{\hat{v}_i^{(k+1)}}{\sqrt{\hat{w}_i^{(k+1)} + \delta}}, \quad i = 1, \dots, r$$

9 $k \leftarrow k + 1$


$$f(x_1, x_2) = x_1^2(x_1 + x_2), \quad \nabla f(2, 3) = (24, 4).$$

$$f(x_1, x_2) = x_1^2(x_1 + x_2), \quad \nabla f(2, 3) = (24, 4).$$




Computation of gradients: Automatic Differentiation

$$f(x_1, x_2) = x_1^2(x_1 + x_2), \quad \nabla f(2, 3) = (24, 4).$$



Backpropagation



Important parameters to keep in mind



Important parameters to keep in mind



- m : number of free parameters

Important parameters to keep in mind

- m : number of free parameters
- n : size of the training dataset

Important parameters to keep in mind

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps

Important parameters to keep in mind

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Important parameters to keep in mind



- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Typically, $m, n, t \rightarrow \infty$ and $d \gg 1$.

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Typically, $m, n, t \rightarrow \infty$ and $d \gg 1$.

Examples where d is large include:

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Typically, $m, n, t \rightarrow \infty$ and $d \gg 1$.

Examples where d is large include:

- radioactive transport equation ($d \geq 5$)

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Typically, $m, n, t \rightarrow \infty$ and $d \gg 1$.

Examples where d is large include:

- radioactive transport equation ($d \geq 5$)
- Boltzmann kinetic equations ($d = 6$)

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Typically, $m, n, t \rightarrow \infty$ and $d \gg 1$.

Examples where d is large include:

- radioactive transport equation ($d \geq 5$)
- Boltzmann kinetic equations ($d = 6$)
- nonlinear Schrödinger equation in the quantum many-body problem ($d \gg 1$)

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Typically, $m, n, t \rightarrow \infty$ and $d \gg 1$.

Examples where d is large include:

- radioactive transport equation ($d \geq 5$)
- Boltzmann kinetic equations ($d = 6$)
- nonlinear Schrödinger equation in the quantum many-body problem ($d \gg 1$)
- parameter-dependent (random) PDEs

- m : number of free parameters
- n : size of the training dataset
- t : number of training steps
- d : input dimension

Typically, $m, n, t \rightarrow \infty$ and $d \gg 1$.

Examples where d is large include:

- radioactive transport equation ($d \geq 5$)
- Boltzmann kinetic equations ($d = 6$)
- nonlinear Schrödinger equation in the quantum many-body problem ($d \gg 1$)
- parameter-dependent (random) PDEs
- nonlinear Black- Scholes equation for pricing derivatives

Deep Learning opens a door to deal with real-world problems



More situations that lead to very large d:

- turbulence modeling,
- plasticity models,
- multiscale,
- multiphysics,
- etc.

Deep Learning opens a door to deal with real-world problems



More situations that lead to very large d :

- turbulence modeling,
- plasticity models,
- multiscale,
- multiphysics,
- etc.



*The heart of the matter for the difficulties described above is our limited ability to handle functions of many variables, and this is exactly where **machine learning** can make a difference.*

Weinan E. The dawning of a new era in applied mathematics , Notice of the AMS, 2021.

<https://web.math.princeton.edu/~weinan/>

Deep Learning opens a door to deal with real-world problems



More situations that lead to very large d :

- turbulence modeling,
- plasticity models,
- multiscale,
- multiphysics,
- etc.



*The heart of the matter for the difficulties described above is our limited ability to handle functions of many variables, and this is exactly where **machine learning** can make a difference.*

Weinan E. The dawning of a new era in applied mathematics , Notice of the AMS, 2021.

<https://web.math.princeton.edu/~weinan/>

Machine learning is a promising tool to deal with **high-dimensional** problems

The methods of this course are (very) easy to implement for a large class of mathematical models. Therefore, they may be useful to:

- test new models very quickly

The methods of this course are (very) easy to implement for a large class of mathematical models. Therefore, they may be useful to:

- test new models very quickly
- incorporate experimental data in the models

The methods of this course are (very) easy to implement for a large class of mathematical models. Therefore, they may be useful to:

- test new models very quickly
- incorporate experimental data in the models
- eventhough predictions maybe not so accurate, they may be used as an initial guess for a more accurate model

Part II

Physics Informed Neural Networks (PINNs)

Goals

Goals

- Study the basis of PINNs algorithm for solving a variety of PDEs-based mathematical models

Goals

- Study the basis of PINNs algorithm for solving a variety of PDEs-based mathematical models
- Gain some computational skills on PINNs -based software

Goals

- Study the basis of PINNs algorithm for solving a variety of PDEs-based mathematical models
- Gain some computational skills on PINNs -based software

Goals

- Study the basis of PINNs algorithm for solving a variety of PDEs-based mathematical models
- Gain some computational skills on PINNs -based software

References



Raisi, M., Perdikaris, P. and Karniadakis, G: *Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations*, *J. Comput. Physics* **378**, 686-707, 2019.



S Wang, S Sankaran, H Wang, P Perdikaris: *An expert's guide to training physics-informed neural networks*. *arXiv:2308.08468*, 178, 2023.



Lu, L., Meng, X., Mao, Z. and Karniadakis, G.: *DeepXDE: A Deep Learning Library for Solving Differential Equations*, *SIAM Review*, **63** (1), 208-228, 2021.



García-Cervera, C.J.; Kessler, M.; Periago, F.: *Control of partial differential equations via physics-informed neural networks*. *J. Optim. Theory Appl.* **196** (2), 391-414, 2023.

A toy model: null control of the wave equation


$$\begin{cases} y_{tt} - \Delta y = 0, & \text{in } Q_T \\ y(x, 0) = y^0(x), & \text{in } \Omega \\ y_t(x, 0) = y^1(x) & \text{in } \Omega \\ y(x, t) = 0, & \text{on } \Gamma_D \times (0, T) \\ y(x, t) = u(x, t) & \text{on } \Gamma_C \times (0, T) \end{cases}$$

Goal: Compute $u(x, t)$ such that

$$y(x, T) = y_t(x, T) = 0 \quad x \in \Omega.$$

Step 1: Neural network

A surrogate $\hat{y}(x, t; \theta)$ of the state variable $y(x, t)$ is constructed as



$$\left\{ \begin{array}{ll} \text{input layer:} & \mathcal{N}^0(\mathbf{x}) = \mathbf{x} = (x, t) \in \mathbb{R}^{d+1} \\ \text{hidden layers:} & \mathcal{N}^\ell(\mathbf{x}) = \sigma(\mathbf{W}^\ell \mathcal{N}^{\ell-1}(\mathbf{x}) + \mathbf{b}^\ell) \in \mathbb{R}^{N_\ell}, \quad \ell = 1, \dots, L-1 \\ \text{output layer:} & \hat{y}(\mathbf{x}; \theta) = \mathcal{N}^L(\mathbf{x}) = \mathbf{W}^L \mathcal{N}^{L-1}(\mathbf{x}) + \mathbf{b}^L \in \mathbb{R} \end{array} \right.$$

- $\mathcal{N}^\ell : \mathbb{R}^{d_{in}} \rightarrow \mathbb{R}^{d_{out}}$ is the ℓ layer with N_ℓ neurons,
- $\mathbf{W}^\ell \in \mathbb{R}^{N_\ell \times N_{\ell-1}}$ and $\mathbf{b}^\ell \in \mathbb{R}^{N_\ell}$ are, respectively, the weights and biases so that $\theta = \{\mathbf{W}^\ell, \mathbf{b}^\ell\}_{1 \leq \ell \leq L}$ are the parameters of the neural network, and
- σ is an activation function, e.g. $\sigma(s) = \tanh(s)$

Step 2: Training dataset




Figure: Illustration of a training dataset (based on Sobol points) in the domain $Q_2 = (0, 1) \times (0, 2)$. Interior points are marked with circles and boundary points in blue color. $\mathbf{x}_j = (x_j, t_j)$ are the features.

Step 3: Loss function. Labels equal zero

$$\mathcal{L}_{\text{int}}(\boldsymbol{\theta}; \mathcal{T}_{\text{int}}) = \sum_{j=1}^{N_{\text{int}}} w_{j,\text{int}} |\hat{y}_{\text{tt}}(\mathbf{x}_j; \boldsymbol{\theta}) - \Delta \hat{y}(\mathbf{x}_j; \boldsymbol{\theta})|^2, \quad \mathbf{x}_j \in \mathcal{T}_{\text{int}}$$

$$\mathcal{L}_{\Gamma_D}(\boldsymbol{\theta}; \mathcal{T}_{\Gamma_D}) = \sum_{j=1}^{N_b} w_{j,b} |\hat{y}(\mathbf{x}_j; \boldsymbol{\theta})|^2, \quad \mathbf{x}_j \in \mathcal{T}_{\Gamma_D}$$

$$\mathcal{L}_{t=0}^{\text{pos}}(\boldsymbol{\theta}; \mathcal{T}_{t=0}) = \sum_{j=1}^{N_0} w_{j,0} |\hat{y}(\mathbf{x}_j; \boldsymbol{\theta}) - y^0(\mathbf{x}_j)|^2, \quad \mathbf{x}_j \in \mathcal{T}_{t=0}$$

$$\mathcal{L}_{t=0}^{\text{vel}}(\boldsymbol{\theta}; \mathcal{T}_{t=0}) = \sum_{j=1}^{N_0} w_{j,0} |\hat{y}_t(\mathbf{x}_j; \boldsymbol{\theta}) - y^1(\mathbf{x}_j)|^2, \quad \mathbf{x}_j \in \mathcal{T}_{t=0}$$

$$\mathcal{L}_{t=T}^{\text{pos}}(\boldsymbol{\theta}; \mathcal{T}_{t=T}) = \sum_{j=1}^{N_T} w_{j,T} |\hat{y}(\mathbf{x}_j; \boldsymbol{\theta})|^2, \quad \mathbf{x}_j \in \mathcal{T}_{t=T}$$

$$\mathcal{L}_{t=T}^{\text{vel}}(\boldsymbol{\theta}; \mathcal{T}_{t=T}) = \sum_{j=1}^{N_T} w_{j,T} |\hat{y}_t(\mathbf{x}_j; \boldsymbol{\theta})|^2, \quad \mathbf{x}_j \in \mathcal{T}_{t=T},$$

where $w_{j,\text{int}}$, $w_{j,b}$, $w_{j,0}$ and $w_{j,T}$ are the weights of suitable quadrature rules.

$$\begin{aligned} \mathcal{L}(\boldsymbol{\theta}; \mathcal{T}) &= \lambda_1 \mathcal{L}_{\text{int}}(\boldsymbol{\theta}; \mathcal{T}_{\text{int}}) \\ &\quad + \lambda_2 \mathcal{L}_{\Gamma_D}(\boldsymbol{\theta}; \mathcal{T}_{\Gamma_D}) \\ &\quad + \lambda_3 \mathcal{L}_{t=0}^{\text{pos}}(\boldsymbol{\theta}; \mathcal{T}_{t=0}) + \lambda_4 \mathcal{L}_{t=0}^{\text{vel}}(\boldsymbol{\theta}; \mathcal{T}_{t=0}) \\ &\quad + \lambda_5 \mathcal{L}_{t=T}^{\text{pos}}(\boldsymbol{\theta}; \mathcal{T}_{t=T}) + \lambda_6 \mathcal{L}_{t=T}^{\text{vel}}(\boldsymbol{\theta}; \mathcal{T}_{t=T}). \end{aligned}$$

Step 4: Training process

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta; \mathcal{T}).$$

The approximation $\hat{u}(t; \theta^*)$ of the control $u(x, t)$ is

$$\hat{u}(x, t; \theta^*) = \hat{y}(x, t; \theta^*), \quad x \in \Gamma_C, \quad 0 \leq t \leq T.$$


Step 4: Training process

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta; \mathcal{T}).$$

The approximation $\hat{u}(t; \theta^*)$ of the control $u(x, t)$ is

$$\hat{u}(x, t; \theta^*) = \hat{y}(x, t; \theta^*), \quad x \in \Gamma_C, \quad 0 \leq t \leq T.$$

To sum up:



Why MLP is a suitable prediction model ?

Why MLP is a suitable prediction model ?

Let us consider the hypothesis space of single-layer neural nets

$$\mathcal{H}_m := \left\{ y_m(\mathbf{x}) := \sum_{i=1}^m a_i \sigma(\omega_i \mathbf{x} + b_i) : \mathbf{x}, \omega_i \in \mathbb{R}^{d+1}, a_i, b_i \in \mathbb{R} \right\}.$$

Why MLP is a suitable prediction model ?

Let us consider the hypothesis space of single-layer neural nets

$$\mathcal{H}_m := \left\{ y_m(\mathbf{x}) := \sum_{i=1}^m a_i \sigma(\omega_i \mathbf{x} + b_i) : \mathbf{x}, \omega_i \in \mathbb{R}^{d+1}, a_i, b_i \in \mathbb{R} \right\}.$$

Theorem (Pinkus Universal Approximation Theorem)

Let $f \in C^k(\mathbb{R}^{d+1})$. Assume that the activation function $\sigma \in C^k(\mathbb{R})$ is not a polynomial. Then, for any compact set $K \subset \mathbb{R}^{d+1}$ and any $\varepsilon > 0$ there exists $m \in \mathbb{N}$ and $y_m \in \mathcal{H}_m$ such that

$$\max_{\mathbf{x} \in K} |D^\ell f(\mathbf{x}) - D^\ell y_m(\mathbf{x})| \leq \varepsilon$$

for all multiindex $\ell \leq k$. Moreover, each $a_i = a_i(f)$ is a continuous linear functional defined on K .



**Pinkus, A.: Approximation theory of the MLP model in neural networks
Acta numer. 8, 143-195, 1999.**

Why is so crucial that the activation function not to be a polynomial?

Why is so crucial that the activation function not to be a polynomial?

If σ is a polynomial of degree m , then $\sigma(\omega x + b)$ is a polynomial of total degree at most m . Thus, $\mathcal{H}_m$ is the space of all polynomials of degree m , which is not dense in $C(K)$, $K \subset \mathbb{R}$ compact.

Why is so crucial that the activation function not to be a polynomial?

If σ is a polynomial of degree m , then $\sigma(\omega x + b)$ is a polynomial of total degree at most m . Thus, $\mathcal{H}_m$ is the space of all polynomials of degree m , which is not dense in $C(K)$, $K \subset \mathbb{R}$ compact.

Conversely, let $\sigma \in C^\infty(\mathbb{R})$. Consider the space

$$\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R}) = \text{span} \{ \sigma(wx + b), \quad w, b \in \mathbb{R} \}.$$

Why is so crucial that the activation function not to be a polynomial?

If σ is a polynomial of degree m , then $\sigma(\omega x + b)$ is a polynomial of total degree at most m . Thus, $\mathcal{H}_m$ is the space of all polynomials of degree m , which is not dense in $C(K)$, $K \subset \mathbb{R}$ compact.

Conversely, let $\sigma \in C^\infty(\mathbb{R})$. Consider the space

$$\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R}) = \text{span} \{ \sigma(wx + b), \quad w, b \in \mathbb{R} \}.$$

A classical result in Calculus states that if $\sigma \in C^\infty$ on any open interval and is not a polynomial thereon, then there exists a point b^* in that interval such that $\sigma^{(k)}(b^*) \neq 0$ for all $k = 0, 1, 2, \dots$.

Why is so crucial that the activation function not to be a polynomial?

If σ is a polynomial of degree m , then $\sigma(\omega x + b)$ is a polynomial of total degree at most m . Thus, $\mathcal{H}_m$ is the space of all polynomials of degree m , which is not dense in $C(K)$, $K \subset \mathbb{R}$ compact.

Conversely, let $\sigma \in C^\infty(\mathbb{R})$. Consider the space

$$\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R}) = \text{span} \{ \sigma(wx + b), \quad w, b \in \mathbb{R} \}.$$

A classical result in Calculus states that if $\sigma \in C^\infty$ on any open interval and is not a polynomial thereon, then there exists a point b^* in that interval such that $\sigma^{(k)}(b^*) \neq 0$ for all $k = 0, 1, 2, \dots$. For $h \neq 0$,

$$\frac{\sigma((w+h)x + b^*) - \sigma(wx + b^*)}{h} \in \mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})$$

and so

$$\frac{d}{dw} \sigma(wx + b^*)|_{w=0} = x\sigma'(b^*) \in \overline{\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})}.$$

Why is so crucial that the activation function not to be a polynomial?

If σ is a polynomial of degree m , then $\sigma(\omega x + b)$ is a polynomial of total degree at most m . Thus, $\mathcal{H}_m$ is the space of all polynomials of degree m , which is not dense in $C(K)$, $K \subset \mathbb{R}$ compact.

Conversely, let $\sigma \in C^\infty(\mathbb{R})$. Consider the space

$$\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R}) = \text{span} \{ \sigma(\omega x + b), \quad \omega, b \in \mathbb{R} \}.$$

A classical result in Calculus states that if $\sigma \in C^\infty$ on any open interval and is not a polynomial thereon, then there exists a point b^* in that interval such that $\sigma^{(k)}(b^*) \neq 0$ for all $k = 0, 1, 2, \dots$. For $h \neq 0$,

$$\frac{\sigma((\omega + h)x + b^*) - \sigma(\omega x + b^*)}{h} \in \mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})$$

and so

$$\frac{d}{d\omega} \sigma(\omega x + b^*)|_{\omega=0} = x \sigma'(b^*) \in \overline{\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})}.$$

Analogously,

$$\frac{d^k}{d\omega^k} \sigma(\omega x + b^*)|_{\omega=0} = x^k \sigma^{(k)}(b^*) \in \overline{\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})}.$$

Why is so crucial that the activation function not to be a polynomial?

If σ is a polynomial of degree m , then $\sigma(\omega x + b)$ is a polynomial of total degree at most m . Thus, $\mathcal{H}_m$ is the space of all polynomials of degree m , which is not dense in $C(K)$, $K \subset \mathbb{R}$ compact.

Conversely, let $\sigma \in C^\infty(\mathbb{R})$. Consider the space

$$\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R}) = \text{span} \{ \sigma(\omega x + b), \quad \omega, b \in \mathbb{R} \}.$$

A classical result in Calculus states that if $\sigma \in C^\infty$ on any open interval and is not a polynomial thereon, then there exists a point b^* in that interval such that $\sigma^{(k)}(b^*) \neq 0$ for all $k = 0, 1, 2, \dots$. For $h \neq 0$,

$$\frac{\sigma((\omega + h)x + b^*) - \sigma(\omega x + b^*)}{h} \in \mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})$$

and so

$$\frac{d}{d\omega} \sigma(\omega x + b^*)|_{\omega=0} = x \sigma'(b^*) \in \overline{\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})}.$$

Analogously,

$$\frac{d^k}{d\omega^k} \sigma(\omega x + b^*)|_{\omega=0} = x^k \sigma^{(k)}(b^*) \in \overline{\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})}.$$

Since $\sigma^{(k)}(b^*) \neq 0$ for all k , then $\overline{\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})}$ contains all polynomials. By Weierstrass theorem, $\mathcal{N}(\sigma; \mathbb{R}, \mathbb{R})$ is dense in $C(K)$ for any compact set $K \subset \mathbb{R}$.

Estimates on generalization error for the null control of the wave equation

Estimates on generalization error for the null control of the wave equation

Training error

$$\mathcal{E}_{\text{train}} := \mathcal{E}_{\text{train, int}} + \mathcal{E}_{\text{train, boundary}} + \mathcal{E}_{\text{train, initialpos}} + \mathcal{E}_{\text{train, initialvel}} \\ + \mathcal{E}_{\text{train, finalpos}} + \mathcal{E}_{\text{train, finalvel}},$$

$$\left\{ \begin{array}{ll} \mathcal{E}_{\text{train, int}} &= (\mathcal{L}_{\text{int}}(\theta^*; \mathcal{T}_{\text{int}}))^{1/2} \\ \mathcal{E}_{\text{train, boundary}} &= (\mathcal{L}_{\Gamma_D}(\theta^*; \mathcal{T}_{\Gamma_D}))^{1/2} \\ \mathcal{E}_{\text{train, initialpos}} &= (\mathcal{L}_{t=0}^{\text{pos}}(\theta^*; \mathcal{T}_{t=0}))^{1/2} \\ \mathcal{E}_{\text{train, initialvel}} &= (\mathcal{L}_{t=0}^{\text{vel}}(\theta^*; \mathcal{T}_{t=0}))^{1/2} \\ \mathcal{E}_{\text{train, finalpos}} &= (\mathcal{L}_{t=T}^{\text{pos}}(\theta^*; \mathcal{T}_{t=T}))^{1/2} \\ \mathcal{E}_{\text{train, finalvel}} &= (\mathcal{L}_{t=T}^{\text{vel}}(\theta^*; \mathcal{T}_{t=T}))^{1/2}, \end{array} \right.$$

Estimates on generalization error for the null control of the wave equation

Training error

$$\mathcal{E}_{\text{train}} := \mathcal{E}_{\text{train, int}} + \mathcal{E}_{\text{train, boundary}} + \mathcal{E}_{\text{train, initialpos}} + \mathcal{E}_{\text{train, initialvel}} + \mathcal{E}_{\text{train, finalpos}} + \mathcal{E}_{\text{train, finalvel}},$$

$$\left\{ \begin{array}{ll} \mathcal{E}_{\text{train, int}} &= (\mathcal{L}_{\text{int}}(\theta^*; \mathcal{T}_{\text{int}}))^{1/2} \\ \mathcal{E}_{\text{train, boundary}} &= (\mathcal{L}_{\Gamma_D}(\theta^*; \mathcal{T}_{\Gamma_D}))^{1/2} \\ \mathcal{E}_{\text{train, initialpos}} &= (\mathcal{L}_{t=0}^{\text{pos}}(\theta^*; \mathcal{T}_{t=0}))^{1/2} \\ \mathcal{E}_{\text{train, initialvel}} &= (\mathcal{L}_{t=0}^{\text{vel}}(\theta^*; \mathcal{T}_{t=0}))^{1/2} \\ \mathcal{E}_{\text{train, finalpos}} &= (\mathcal{L}_{t=T}^{\text{pos}}(\theta^*; \mathcal{T}_{t=T}))^{1/2} \\ \mathcal{E}_{\text{train, finalvel}} &= (\mathcal{L}_{t=T}^{\text{vel}}(\theta^*; \mathcal{T}_{t=T}))^{1/2}, \end{array} \right.$$

Generalization error for control and state

$$\left\{ \begin{array}{l} \mathcal{E}_{\text{gener}}(u) := \|u - \hat{u}\|_{L^2(\Gamma_C; (0, T))} \\ \mathcal{E}_{\text{gener}}(y) := \|y - \hat{y}\|_{C(0, T; L^2(\Omega)) \cap C^1(0, T; H^{-1}(\Omega))} \end{array} \right.$$

Theorem (Estimates on generalization error)

Assume that both $y, \hat{y} \in C^2(\overline{Q_T})$. Then

$$\begin{aligned} \mathcal{E}_{gener}(u) \leq & C \left(\mathcal{E}_{train, int} + C_{int}^{1/2} N_{int}^{-\alpha_{int}/2} \right. \\ & + \mathcal{E}_{train, boundary} + C_{qb}^{1/2} N_b^{-\alpha_b/2} \\ & + \mathcal{E}_{train, initialpos} + C_{qip}^{1/2} N_0^{-\alpha_{ip}/2} \\ & + \mathcal{E}_{train, initialvel} + C_{qiv}^{1/2} N_0^{-\alpha_{iv}/2} \\ & + \mathcal{E}_{train, finalpos} + C_{qfp}^{1/2} N_T^{-\alpha_{fp}/2} \\ & \left. + \mathcal{E}_{train, finalvel} + C_{fv}^{1/2} N_T^{-\alpha_{fv}/2} \right), \end{aligned}$$

where $C = C(\Omega, T)$, and consequently $C = C(d)$ also depends on the spatial dimension d . A similar estimate holds for the state variable.

Moreover, training errors converge to zero as the size of the NN and the number of training points go to infinity.



García-Cervera, C., Kessler, M., Periago, F.: Control of Partial Differential Equations via Physics-Informed Neural Networks **J. Optim. Th. Appl.** **196**, 391–414, 2023.

- **Data normalization.** It is important to ensure that the target output variables vary within a reasonable range. One way to achieve this is through non-dimensionalization of the PDE.

- **Data normalization.** It is important to ensure that the target output variables vary within a reasonable range. One way to achieve this is through non-dimensionalization of the PDE.
- **Network Architecture.** It is recommended to use Multi-layer Perceptrons (MLPs) with depth and width ranging from 3 to 6, and 128 to 512, respectively. As for the activation function, the hyperbolic tangent. For initialization of the network parameters, a Glorot uniform distribution.

- **Data normalization.** It is important to ensure that the target output variables vary within a reasonable range. One way to achieve this is through non-dimensionalization of the PDE.
- **Network Architecture.** It is recommended to use Multi-layer Perceptrons (MLPs) with depth and width ranging from 3 to 6, and 128 to 512, respectively. As for the activation function, the hyperbolic tangent. For initialization of the network parameters, a Glorot uniform distribution.
- **Loss balancing.** $\mathcal{L}(\theta) = \lambda_{ic}\mathcal{L}_{ic}(\theta) + \lambda_{bc}\mathcal{L}_{bc}(\theta) + \lambda_r\mathcal{L}_r(\theta)$

- **Data normalization.** It is important to ensure that the target output variables vary within a reasonable range. One way to achieve this is through non-dimensionalization of the PDE.
- **Network Architecture.** It is recommended to use Multi-layer Perceptrons (MLPs) with depth and width ranging from 3 to 6, and 128 to 512, respectively. As for the activation function, the hyperbolic tangent. For initialization of the network parameters, a Glorot uniform distribution.
- **Loss balancing.** $\mathcal{L}(\theta) = \lambda_{ic}\mathcal{L}_{ic}(\theta) + \lambda_{bc}\mathcal{L}_{bc}(\theta) + \lambda_r\mathcal{L}_r(\theta)$

1 Compute

$$\left\{ \begin{array}{l} \hat{\lambda}_{ic} = \frac{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_r(\theta)\|}{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\|} \\ \hat{\lambda}_{bc} = \frac{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_r(\theta)\|}{\|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\|} \\ \hat{\lambda}_r = \frac{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_r(\theta)\|}{\|\nabla_{\theta}\mathcal{L}_r(\theta)\|} \end{array} \right.$$

- **Data normalization.** It is important to ensure that the target output variables vary within a reasonable range. One way to achieve this is through non-dimensionalization of the PDE.
- **Network Architecture.** It is recommended to use Multi-layer Perceptrons (MLPs) with depth and width ranging from 3 to 6, and 128 to 512, respectively. As for the activation function, the hyperbolic tangent. For initialization of the network parameters, a Glorot uniform distribution.
- **Loss balancing.** $\mathcal{L}(\theta) = \lambda_{ic}\mathcal{L}_{ic}(\theta) + \lambda_{bc}\mathcal{L}_{bc}(\theta) + \lambda_r\mathcal{L}_r(\theta)$

1 Compute

$$\left\{ \begin{array}{l} \hat{\lambda}_{ic} = \frac{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_r(\theta)\|}{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\|} \\ \hat{\lambda}_{bc} = \frac{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_r(\theta)\|}{\|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\|} \\ \hat{\lambda}_r = \frac{\|\nabla_{\theta}\mathcal{L}_{ic}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_{bc}(\theta)\| + \|\nabla_{\theta}\mathcal{L}_r(\theta)\|}{\|\nabla_{\theta}\mathcal{L}_r(\theta)\|} \end{array} \right.$$

2 Update the weights $\lambda = (\lambda_{ic}, \lambda_{bc}, \lambda_r)$ by using a moving average

$$\lambda_{\text{new}} = \alpha\lambda_{\text{old}} + (1 - \alpha)\hat{\lambda}_{\text{new}}$$

- **Random weight factorization** may improve the performance of PINNs:

$$w^{(k,\ell)} = s^{(k,\ell)} v^{(k,\ell)},$$

with $s^{(k,\ell)}$ is a trainable scale factor assigned to each individual neuron.



S. Wang, et al.: *Random weight factorization improves the training of continuous neural representations*. [arXiv:2210.01274](https://arxiv.org/abs/2210.01274), 2022.

- **Random weight factorization** may improve the performance of PINNs:

$$w^{(k,\ell)} = s^{(k,\ell)} v^{(k,\ell)},$$

with $s^{(k,\ell)}$ is a trainable scale factor assigned to each individual neuron.



S. Wang, et al.: *Random weight factorization improves the training of continuous neural representations*. [arXiv:2210.01274, 2022](#).

- **Optimization algorithm.** The Adaptive with Moment (Adam) algorithm is the most widely used to minimise the loss function. To speed up convergence near a local minimum, it is convenient to combine Adam (say, first 20000 iterations) with a quasi-Newton method like L-BFGS.

- PINN has difficulty in approximating functions that have **steep gradients**. A Residual-based Adaptive Refinement algorithm has been proposed which add more residual points in the locations where the PDE residual is large.



Lu, L., Meng, X., Mao, Z. and Karniadakis, G.: *DeepXDE: A Deep Learning Library for Solving Differential Equations*, **SIAM Review**, 63 (1), 208-228, 2021.

- PINN has difficulty in approximating functions that have **steep gradients**. A Residual-based Adaptive Refinement algorithm has been proposed which add more residual points in the locations where the PDE residual is large.

 Lu, L., Meng, X., Mao, Z. and Karniadakis, G.: *DeepXDE: A Deep Learning Library for Solving Differential Equations*, **SIAM Review**, 63 (1), 208-228, 2021.

- **MLPs are biased towards learning low frequencies functions.** This can be mitigated by using a random Fourier feature embedding like

$$\gamma(x) = (\cos(Bx), \sin(Bx)),$$

where B_{ij} are sampled from Gaussians. This maps input coordinates into high frequency signals before passing through the network.

 M. Tancik, et al.: *Fourier features let networks learn high frequency functions in low dimensional domains*. **arXiv:2006.10739**, 2020.

- PINN has difficulty in approximating functions that have **steep gradients**. A Residual-based Adaptive Refinement algorithm has been proposed which add more residual points in the locations where the PDE residual is large.

 Lu, L., Meng, X., Mao, Z. and Karniadakis, G.: *DeepXDE: A Deep Learning Library for Solving Differential Equations*, **SIAM Review**, **63** (1), 208-228, 2021.

- **MLPs are biased towards learning low frequencies functions**. This can be mitigated by using a random Fourier feature embedding like

$$\gamma(x) = (\cos(Bx), \sin(Bx)),$$

where B_{ij} are sampled from Gaussians. This maps input coordinates into high frequency signals before passing through the network.

 M. Tancik, et al.: *Fourier features let networks learn high frequency functions in low dimensional domains*. **arXiv:2006.10739**, 2020.

- **PINNs may violate temporal causality** when solving time-dependent PDEs. We should partition the temporal domain into M equal sequential segments and introduce $\mathcal{L}_r^i(\theta)$ to denote the PDE residual loss within the i -th segment. The PDE residual becomes $\mathcal{L}_r(\theta) = \sum_{i=1}^M \lambda_i \mathcal{L}_r^i(\theta)$.

 S Wang, S Sankaran, H Wang, P Perdikaris: *An expert's guide to training physics-informed neural networks*. **ArXiv:2308.08468**, 2023.

Part III

Deep Operator Network (DeepONet)

Goals

Goals

- Given two function spaces X and Y , and an operator

$$\mathcal{G} : X \rightarrow Y$$

the main goal is to learn (approximate) the operator $\mathcal{G}$ from a dataset.

Goals

- Given two function spaces X and Y , and an operator

$$\mathcal{G} : X \rightarrow Y$$

the main goal is to learn (approximate) the operator $\mathcal{G}$ from a dataset.

- Gain some computational skills on DeepONet -based software

Goals

- Given two function spaces X and Y , and an operator

$$\mathcal{G} : X \rightarrow Y$$

the main goal is to learn (approximate) the operator $\mathcal{G}$ from a dataset.

- Gain some computational skills on DeepONet -based software

References

-  Lu, L., Jin, P., Pang, G., Zhang, Z., Karniadakis, G.E.: *Learning nonlinear operators via deeponet based on the universal approximation theorem of operators*. **Nature Machine Intelligence** 3(3), 218-229, 2021.
-  Lanthaler, S., Mishra, S., Karniadakis, G.E.: *Error estimates for DeepONets: a deep learning framework in infinite dimensions*. **Trans. Math. Appl.** 6(1), 1–144, 2022.
-  García-Cervera, C.J.; Kessler, M.; Pedregal, P.; Periago, F.: *Universal Approximation of Set-Valued Maps and Application to Control*. **Submitted**, 2025.

For the sake of clarity, we focus on the control problem

$$\left\{ \begin{array}{ll} y_{tt} - y_{xx} = 0, & \text{in } (0, 1) \times (0, 2) \\ y(x, 0) = y^0(x), & \text{on } (0, 1) \\ y_t(x, 0) = y^1(x) & \text{on } (0, 1) \\ y(0, t) = 0, & \text{on } (0, 2) \\ y(1, t) = u(t) & \text{on } (0, 2) \\ y(x, 2) = y_t(x, 2) = 0, & \text{on } (0, 1). \end{array} \right.$$

For the sake of clarity, we focus on the control problem

$$\left\{ \begin{array}{ll} y_{tt} - y_{xx} = 0, & \text{in } (0, 1) \times (0, 2) \\ y(x, 0) = y^0(x), & \text{on } (0, 1) \\ y_t(x, 0) = y^1(x) & \text{on } (0, 1) \\ y(0, t) = 0, & \text{on } (0, 2) \\ y(1, t) = u(t) & \text{on } (0, 2) \\ y(x, 2) = y_t(x, 2) = 0, & \text{on } (0, 1). \end{array} \right.$$

Goal: approximate the controllability mapping

$$\begin{aligned} \mathcal{G} : \quad L^2(0, 1) \times H^{-1}(0, 1) &\rightarrow L^2(0, 2) \\ (y^0, y^1) &\mapsto \mathcal{G}(y^0, y^1) := u \end{aligned}$$

where u is the **unique** control of minimal L^2 -norm.

For the sake of clarity, we focus on the control problem

$$\left\{ \begin{array}{ll} y_{tt} - y_{xx} = 0, & \text{in } (0, 1) \times (0, 2) \\ y(x, 0) = y^0(x), & \text{on } (0, 1) \\ y_t(x, 0) = y^1(x) & \text{on } (0, 1) \\ y(0, t) = 0, & \text{on } (0, 2) \\ y(1, t) = u(t) & \text{on } (0, 2) \\ y(x, 2) = y_t(x, 2) = 0, & \text{on } (0, 1). \end{array} \right.$$

Goal: approximate the controllability mapping

$$\begin{aligned} \mathcal{G} : \quad L^2(0, 1) \times H^{-1}(0, 1) &\rightarrow L^2(0, 2) \\ (y^0, y^1) &\mapsto \mathcal{G}(y^0, y^1) := u \end{aligned}$$

where u is the **unique** control of minimal L^2 -norm.

The operator $\mathcal{G}$ is well-defined (uniqueness of the control), linear and continuous. Continuity is a consequence of the observability inequality

$$\|u\|_{L^2(0,2)} \leq C \left(\|y^0\|_{L^2(0,1)} + \|y^1\|_{H^{-1}(0,1)} \right)$$

Thus, $\mathcal{G}$ is Lipschitz continuous.

■ Dataset

We fix a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$. The information of each selected continuous initial datum (y^0, y^1) is encoded in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

■ Dataset

We fix a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$. The information of each selected continuous initial datum (y^0, y^1) is encoded in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

We also take $t \in [0, 2]$.

■ Dataset

We fix a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$. The information of each selected continuous initial datum (y^0, y^1) is encoded in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

We also take $t \in [0, 2]$. Putting all together,

$$\{(y_\ell^{\text{initial}}; t_\ell), \quad 1 \leq \ell \leq N\}$$

■ Dataset

We fix a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$. The information of each selected continuous initial datum (y^0, y^1) is encoded in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

We also take $t \in [0, 2]$. Putting all together,

$$\{(y_\ell^{\text{initial}}; t_\ell), \quad 1 \leq \ell \leq N\}$$

The corresponding **labels** are

$$\{u_\ell = u(y_\ell^{\text{initial}}; t_\ell), \quad 1 \leq \ell \leq N\},$$

the control at time t_ℓ associated with the initial datum y_ℓ^{initial} .

■ Dataset

We fix a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$. The information of each selected continuous initial datum (y^0, y^1) is encoded in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

We also take $t \in [0, 2]$. Putting all together,

$$\{(y_\ell^{\text{initial}}; t_\ell), \quad 1 \leq \ell \leq N\}$$

The corresponding **labels** are

$$\{u_\ell = u(y_\ell^{\text{initial}}; t_\ell), \quad 1 \leq \ell \leq N\},$$

the control at time t_ℓ associated with the initial datum y_ℓ^{initial} .

■ Hypothesis space: the neural network

We will use the so-called **DeepONet**, which takes the form

$$\mathcal{N}(\theta; (y^{\text{initial}}(x_j); t)) := \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k y^{\text{initial}}(x_j) + \theta_i^k \right) \cdot \sigma(w_k \cdot t + \eta_k)$$

where $\theta = (c_i^k, \xi_{ij}^k, \theta_i^k, w_k, \eta_k)$ is the set of parameters of the net.

■ Dataset

We fix a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$. The information of each selected continuous initial datum (y^0, y^1) is encoded in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

We also take $t \in [0, 2]$. Putting all together,

$$\{(y_\ell^{\text{initial}}; t_\ell), \quad 1 \leq \ell \leq N\}$$

The corresponding **labels** are

$$\{u_\ell = u(y_\ell^{\text{initial}}; t_\ell), \quad 1 \leq \ell \leq N\},$$

the control at time t_ℓ associated with the initial datum y_ℓ^{initial} .

■ Hypothesis space: the neural network

We will use the so-called **DeepONet**, which takes the form

$$\mathcal{N}(\theta; (y^{\text{initial}}(x_j); t)) := \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k y^{\text{initial}}(x_j) + \theta_i^k \right) \cdot \sigma(w_k \cdot t + \eta_k)$$


where $\theta = (c_i^k, \xi_{ij}^k, \theta_i^k, w_k, \eta_k)$ is the set of parameters of the net.

■ Loss function: Mean Squared Error (MSE)

$$\text{MSE}(\theta) = \frac{1}{N} \sum_{\ell=1}^N |\mathcal{N}(\theta; (y_\ell^{\text{initial}}; t_\ell)) - u_\ell|^2$$

DeepONet's architecture

$$\mathcal{N}(\theta; (y^{\text{initial}}(x_j); t)) := \underbrace{\sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k y^{\text{initial}}(x_j) + \theta_i^k \right)}_{\beta_k} \cdot \underbrace{\sigma(w_k \cdot t + \eta_k)}_{\tau_k}$$



From a different perspective, the approximation scheme may be decomposed as

$$\begin{array}{ccc} X & \xrightarrow{\mathcal{G}} & Y \\ \downarrow \mathcal{E} & & \uparrow \mathcal{R} \\ \mathbb{R}^{2m} & \xrightarrow{\mathcal{A}} & \mathbb{R}^p \end{array}$$

where:

From a different perspective, the approximation scheme may be decomposed as

$$\begin{array}{ccc} X & \xrightarrow{\mathcal{G}} & Y \\ \downarrow \mathcal{E} & & \uparrow \mathcal{R} \\ \mathbb{R}^{2m} & \xrightarrow{\mathcal{A}} & \mathbb{R}^p \end{array}$$

where:

- $\mathcal{E} : (y^0, y^1) \mapsto y^{initial}$ is an *encoder*,

From a different perspective, the approximation scheme may be decomposed as

$$\begin{array}{ccc} X & \xrightarrow{\mathcal{G}} & Y \\ \downarrow \mathcal{E} & & \uparrow \mathcal{R} \\ \mathbb{R}^{2m} & \xrightarrow{\mathcal{A}} & \mathbb{R}^p \end{array}$$

where:

- $\mathcal{E} : (y^0, y^1) \mapsto y^{initial}$ is an *encoder*,
- $\mathcal{A} : y^{initial} \mapsto \beta_k(y^{initial})$ is a finite dimensional *approximation operator*, and

From a different perspective, the approximation scheme may be decomposed as

$$\begin{array}{ccc} X & \xrightarrow{\mathcal{G}} & Y \\ \downarrow \mathcal{E} & & \uparrow \mathcal{R} \\ \mathbb{R}^{2m} & \xrightarrow{\mathcal{A}} & \mathbb{R}^p \end{array}$$

where:

- $\mathcal{E} : (y^0, y^1) \mapsto y^{initial}$ is an *encoder*,
- $\mathcal{A} : y^{initial} \mapsto \beta_k(y^{initial})$ is a finite dimensional *approximation operator*, and
- $\mathcal{R} : \beta_k(y^{initial}) \mapsto \sum_{k=1}^p \beta_k(y^{initial}) \tau_k(t)$ is a *reconstruction operator*.

Thus,

$$\mathcal{G} \approx \mathcal{R} \circ \mathcal{A} \circ \mathcal{E}.$$

Where does this architecture come from ?

Theorem (Universal Approximation Theorem for Functions)

Suppose that $K \subset \mathbb{R}^d$ is compact, $U \subset C(K)$ is compact, and $\sigma \in C^\infty(\mathbb{R})$ is not a polynomial. Then, for any $\varepsilon > 0$ there exist a positive integer n , real numbers $\theta_i, \omega_i \in \mathbb{R}^n$, independent of $f \in U$, and constants $c_i = c_i(f)$ depending on f , such that

$$|f(x) - \sum_{i=1}^n c_i \sigma(\omega_i \cdot x + \theta_i)| < \varepsilon$$

holds for all $x \in K$ and $f \in U$. Moreover, each $c_i(f)$ is a continuous linear functional defined on U .

Where does this architecture come from ?

Theorem (Universal Approximation Theorem for Functions)

Suppose that $K \subset \mathbb{R}^d$ is compact, $U \subset C(K)$ is compact, and $\sigma \in C^\infty(\mathbb{R})$ is not a polynomial. Then, for any $\varepsilon > 0$ there exist a positive integer n , real numbers $\theta_i, \omega_i \in \mathbb{R}^n$, independent of $f \in U$, and constants $c_i = c_i(f)$ depending on f , such that

$$|f(x) - \sum_{i=1}^n c_i \sigma(\omega_i \cdot x + \theta_i)| < \varepsilon$$

holds for all $x \in K$ and $f \in U$. Moreover, each $c_i(f)$ is a continuous linear functional defined on U .

Theorem (Universal Approximation Theorem for Functionals)

Suppose that σ is not a polynomial, X is a Banach space, $K \subset X$ is a compact set, V is a compact set in $C(K)$, and $f : V \rightarrow \mathbb{R}$ is a continuous functional. Then for any $\varepsilon > 0$, there are a positive integer n, m sensor points $x_1, x_2, \dots, x_m \in K$, and real constants $c_i, \theta_i, \xi_{ij}, 1 \leq i \leq n, 1 \leq j \leq m$, such that

$$|f(y) - \sum_{i=1}^n c_i \sigma \left(\sum_{j=1}^m \xi_{ij} y(x_j) + \theta_i \right)| < \varepsilon, \quad \text{for all } y \in V.$$

Where does this architecture come from ?



if $\mathcal{G}$ is an operator between two function spaces, from the Universal Approximation Theorem for functions we get

$$\mathcal{G}(y)(t) \approx \sum_{k=1}^n c_k (\mathcal{G}(y)) \sigma(\omega_k \cdot t + \theta_k),$$

where $c_k (\mathcal{G}(y))$ are continuous functionals.

Where does this architecture come from ?

if $\mathcal{G}$ is an operator between two function spaces, from the Universal Approximation Theorem for functions we get

$$\mathcal{G}(y)(t) \approx \sum_{k=1}^n c_k (\mathcal{G}(y)) \sigma(\omega_k \cdot t + \theta_k),$$

where $c_k (\mathcal{G}(y))$ are continuous functionals.

Theorem (Universal Approximation Theorem for Operators)

Suppose that σ is not a polynomial, X is a Banach space, $K_1 \subset X$, $K_2 \subset \mathbb{R}^d$ are compact sets, V is a compact set in $C(K_1)$, and $\mathcal{G} : V \rightarrow C(K_2)$ is a continuous operator. Then for any $\varepsilon > 0$, there are a positive integers n, p , m sensor points $x_1, x_2, \dots, x_m \in K_1$, and real constants $c_i^k, \theta_i^k, \xi_{ij}^k, \eta_k$ such that

$$|\mathcal{G}(y)(t) - \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k y(x_j) + \theta_i^k \right) \sigma(\omega_k \cdot t + \zeta_k)| < \varepsilon, \quad \forall y \in V, t \in K_2$$



Chen, T., Chen, H.: Universal approximation to nonlinear operators by neural networks with arbitrary activation functions and its application to dynamical systems. **IEEE Transactions on Neural Networks** 6(4), 911-917, 1995.

Definition (Data for DeepONet approximation)

Assume that $X \hookrightarrow L^2(D)$, and $Y \hookrightarrow L^2(U)$, for some Banach spaces X, Y . The pair $(\mu, \mathcal{G})$ is said to be **data for DeepONet approximation** provided $\mu \in \mathcal{P}_2(X)$ is Borel measurable with finite second moments, there exists a Borel set $A \subset X$, composed of continuous functions, $\mu(A) = 1$, and $\mathcal{G} : X \rightarrow Y$ is a Borel measurable mapping with $\|\mathcal{G}\|_{L^2(\mu)} < \infty$.

Definition (Data for DeepONet approximation)

Assume that $X \hookrightarrow L^2(D)$, and $Y \hookrightarrow L^2(U)$, for some Banach spaces X, Y . The pair $(\mu, \mathcal{G})$ is said to be **data for DeepONet approximation** provided $\mu \in \mathcal{P}_2(X)$ is Borel measurable with finite second moments, there exists a Borel set $A \subset X$, composed of continuous functions, $\mu(A) = 1$, and $\mathcal{G} : X \rightarrow Y$ is a Borel measurable mapping with $\|\mathcal{G}\|_{L^2(\mu)} < \infty$.

We need an **encoding operator** $\mathcal{E} : X \rightarrow \mathbb{R}^m$ to represent functions as vectors

Definition (Data for DeepONet approximation)

Assume that $X \hookrightarrow L^2(D)$, and $Y \hookrightarrow L^2(U)$, for some Banach spaces X, Y . The pair $(\mu, \mathcal{G})$ is said to be **data for DeepONet approximation** provided $\mu \in \mathcal{P}_2(X)$ is Borel measurable with finite second moments, there exists a Borel set $A \subset X$, composed of continuous functions, $\mu(A) = 1$, and $\mathcal{G} : X \rightarrow Y$ is a Borel measurable mapping with $\|\mathcal{G}\|_{L^2(\mu)} < \infty$.

We need an **encoding operator** $\mathcal{E} : X \rightarrow \mathbb{R}^m$ to represent functions as vectors
How to sample continuous functions for the initial data (y^0, y^1) ?

Definition (Data for DeepONet approximation)

Assume that $X \hookrightarrow L^2(D)$, and $Y \hookrightarrow L^2(U)$, for some Banach spaces X, Y . The pair $(\mu, \mathcal{G})$ is said to be **data for DeepONet approximation** provided $\mu \in \mathcal{P}_2(X)$ is Borel measurable with finite second moments, there exists a Borel set $A \subset X$, composed of continuous functions, $\mu(A) = 1$, and $\mathcal{G} : X \rightarrow Y$ is a Borel measurable mapping with $\|\mathcal{G}\|_{L^2(\mu)} < \infty$.

We need an **encoding operator** $\mathcal{E} : X \rightarrow \mathbb{R}^m$ to represent functions as vectors

How to sample continuous functions for the initial data (y^0, y^1) ?

We choose $\mu \in \mathcal{P}(C(0, 1))$ a probability measure whose probability law is given by a Karhunen-Loève expansion

$$\mu \sim \sum_{j=1}^{\infty} \sqrt{\lambda_j} \xi_j(\omega) \varphi_j(x) \quad (3)$$

We take $\xi_j \sim \mathcal{N}(0, 1)$ i.i.d. standard Gaussian variables, and (λ_j, φ_j) are the eigenvalues and normalized eigenfunctions of the operator

$$\mathcal{C}(\phi)(x) = \int_0^1 C(x, x') \phi(x') dx', \quad C(x, x') = \sigma^2 \exp\left(-\frac{|x - x'|^2}{2\ell^2}\right)$$

Definition (Data for DeepONet approximation)

Assume that $X \hookrightarrow L^2(D)$, and $Y \hookrightarrow L^2(U)$, for some Banach spaces X, Y . The pair $(\mu, \mathcal{G})$ is said to be **data for DeepONet approximation** provided $\mu \in \mathcal{P}_2(X)$ is Borel measurable with finite second moments, there exists a Borel set $A \subset X$, composed of continuous functions, $\mu(A) = 1$, and $\mathcal{G} : X \rightarrow Y$ is a Borel measurable mapping with $\|\mathcal{G}\|_{L^2(\mu)} < \infty$.

We need an **encoding operator** $\mathcal{E} : X \rightarrow \mathbb{R}^m$ to represent functions as vectors

How to sample continuous functions for the initial data (y^0, y^1) ?

We choose $\mu \in \mathcal{P}(C(0, 1))$ a probability measure whose probability law is given by a Karhunen-Loève expansion

$$\mu \sim \sum_{j=1}^{\infty} \sqrt{\lambda_j} \xi_j(\omega) \varphi_j(x) \quad (3)$$

We take $\xi_j \sim \mathcal{N}(0, 1)$ i.i.d. standard Gaussian variables, and (λ_j, φ_j) are the eigenvalues and normalized eigenfunctions of the operator

$$\mathcal{C}(\phi)(x) = \int_0^1 C(x, x') \phi(x') dx', \quad C(x, x') = \sigma^2 \exp\left(-\frac{|x - x'|^2}{2\ell^2}\right)$$

Thus, we truncate (3) and sample the Gaussian random variables. Remember that by Mercer's theorem $\{\varphi_j\}$ is an orthonormal basis of $L^2(0, 1)$.

After fixing a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$, the information of each selected continuous initial datum (y^0, y^1) is **encoded** in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

After fixing a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$, the information of each selected continuous initial datum (y^0, y^1) is **encoded** in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

The encoding operator $\mathcal{E} : C(0, 1) \times C(0, 1) \rightarrow \mathbb{R}^m \times \mathbb{R}^m$ maps $(y^0, y^1) \mapsto y^{\text{initial}}$

After fixing a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$, the information of each selected continuous initial datum (y^0, y^1) is **encoded** in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

The encoding operator $\mathcal{E} : C(0, 1) \times C(0, 1) \rightarrow \mathbb{R}^m \times \mathbb{R}^m$ maps $(y^0, y^1) \mapsto y^{\text{initial}}$. We repeat this process for a number of initial conditions (y_j^0, y_j^1) and store the corresponding y_{jk}^{initial} in the rows of the so-called matrix of **features**.

After fixing a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$, the information of each selected continuous initial datum (y^0, y^1) is **encoded** in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

The encoding operator $\mathcal{E} : C(0, 1) \times C(0, 1) \rightarrow \mathbb{R}^m \times \mathbb{R}^m$ maps $(y^0, y^1) \mapsto y^{\text{initial}}$. We repeat this process for a number of initial conditions (y_j^0, y_j^1) and store the corresponding y_{jk}^{initial} in the rows of the so-called matrix of **features**.

The controls $u_j(t) := u_j(t; (y_j^0, y_j^1))$, evaluated at some time t are the so-called vector of **labels**.

After fixing a set of **sensor points** $\{x_1, x_2, \dots, x_m\} \subset [0, 1]$, the information of each selected continuous initial datum (y^0, y^1) is **encoded** in the vector

$$(y^0(x_1), y^0(x_2), \dots, y^0(x_m); y^1(x_1), y^1(x_2), \dots, y^1(x_m)) \equiv y^{\text{initial}}$$

The encoding operator $\mathcal{E} : C(0, 1) \times C(0, 1) \rightarrow \mathbb{R}^m \times \mathbb{R}^m$ maps $(y^0, y^1) \mapsto y^{\text{initial}}$. We repeat this process for a number of initial conditions (y_j^0, y_j^1) and store the corresponding y_{jk}^{initial} in the rows of the so-called matrix of **features**.

The controls $u_j(t) := u_j(t; (y_j^0, y_j^1))$, evaluated at some time t are the so-called vector of **labels**.

Putting all together, the **training dataset** is

$$\{(y_{jk}^{\text{initial}}; u_j(t)), 1 \leq j \leq N, 1 \leq k \leq 2m\},$$

evaluated at a finite selection of times t . Precisely,

$$\begin{bmatrix} y_{1,1}^{\text{initial}} & \cdots & y_{1,2m}^{\text{initial}} & u_1(t_1) \\ \vdots & \vdots & \vdots & \vdots \\ y_{1,1}^{\text{initial}} & \cdots & y_{1,2m}^{\text{initial}} & u_1(t_\ell) \\ \hline \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ \hline y_{N,1}^{\text{initial}} & \cdots & y_{N,2m}^{\text{initial}} & u_N(t_1) \\ \vdots & \vdots & \vdots & \vdots \\ y_{N,1}^{\text{initial}} & \cdots & y_{N,2m}^{\text{initial}} & u_N(t_\ell) \end{bmatrix}$$

Error analysis and the curse of dimensionality



When we talk about error, what are we talking about?

When we talk about error, what are we talking about?

- **Approximation error:** This is the distance between the Hypothesis space and the operator $\mathcal{G}$ to be approximated, i.e., if $\mathcal{F}$ is a fixed space of DeepONets, then

$$\mathcal{N}_{\mathcal{F}} = \arg \min_{\mathcal{N} \in \mathcal{F}} \mathcal{L}(\mathcal{N}) := \int_{L^2(D)} \int_U |\mathcal{G}(y)(t) - \mathcal{N}_{\mathcal{F}}(y)(t)|^2 dt d\mu(y),$$

then **approximation error** is

$$\mathcal{E}_{\text{approx}} := \sqrt{\mathcal{L}(\mathcal{N}_{\mathcal{F}})}.$$

When we talk about error, what are we talking about?

- **Approximation error:** This is the distance between the Hypothesis space and the operator $\mathcal{G}$ to be approximated, i.e., if $\mathcal{F}$ is a fixed space of DeepONets, then

$$\mathcal{N}_{\mathcal{F}} = \arg \min_{\mathcal{N} \in \mathcal{F}} \mathcal{L}(\mathcal{N}) := \int_{L^2(D)} \int_U |\mathcal{G}(y)(t) - \mathcal{N}_{\mathcal{F}}(y)(t)|^2 dt d\mu(y),$$

then **approximation error** is

$$\mathcal{E}_{\text{approx}} := \sqrt{\mathcal{L}(\mathcal{N}_{\mathcal{F}})}.$$

- **Generalization (or estimation) error:** We approximate $\mathcal{N}_{\mathcal{F}}$ by using a specific (training) dataset $\mathcal{T}$, and a **empirical loss**. So, we get

$$\mathcal{N}_{\mathcal{T}} = \arg \min_{\mathcal{N} \in \mathcal{F}} \mathcal{L}_M(\mathcal{N}) := \frac{|U|}{M} \sum_{j=1}^M |\mathcal{G}(y_j)(t_j) - \mathcal{N}(y_j)(t_j)|^2$$

Thus, $\mathcal{E}_{\text{gener}} := (\mathcal{L}(\mathcal{N}_{\mathcal{T}}) - \mathcal{L}(\mathcal{N}_{\mathcal{F}}))^{1/2}$.

When we talk about error, what are we talking about?

- **Approximation error:** This is the distance between the Hypothesis space and the operator $\mathcal{G}$ to be approximated, i.e., if $\mathcal{F}$ is a fixed space of DeepONets, then

$$\mathcal{N}_{\mathcal{F}} = \arg \min_{\mathcal{N} \in \mathcal{F}} \mathcal{L}(\mathcal{N}) := \int_{L^2(D)} \int_U |\mathcal{G}(y)(t) - \mathcal{N}_{\mathcal{F}}(y)(t)|^2 dt d\mu(y),$$

then **approximation error** is

$$\mathcal{E}_{\text{approx}} := \sqrt{\mathcal{L}(\mathcal{N}_{\mathcal{F}})}.$$

- **Generalization (or estimation) error:** We approximate $\mathcal{N}_{\mathcal{F}}$ by using a specific (training) dataset $\mathcal{T}$, and a **empirical loss**. So, we get

$$\mathcal{N}_{\mathcal{T}} = \arg \min_{\mathcal{N} \in \mathcal{F}} \mathcal{L}_M(\mathcal{N}) := \frac{|U|}{M} \sum_{j=1}^M |\mathcal{G}(y_j)(t_j) - \mathcal{N}(y_j)(t_j)|^2$$

Thus, $\mathcal{E}_{\text{gener}} := (\mathcal{L}(\mathcal{N}_{\mathcal{T}}) - \mathcal{L}(\mathcal{N}_{\mathcal{F}}))^{1/2}$.

- **Optimization error:** The empirical loss is highly nonlinear, non-convex so that we compute a local minimum $\mathcal{N}_M$ of $\mathcal{L}_M$. Optimization error is then

$$\mathcal{E}_{\text{optim}} := \|\mathcal{N}_M - \mathcal{N}_{\mathcal{T}}\|^{1/2}.$$

$$\mathcal{E}_{\text{total}} = \mathcal{E}_{\text{approx}} + \mathcal{E}_{\text{gener}} + \mathcal{E}_{\text{optim}}$$

Definition (Curse of dimensionality)

For a given $\varepsilon > 0$, let $\mathcal{N}_\varepsilon$ be a DeepONet such that $\mathcal{E}_{\text{approx}} < \varepsilon$, and

$$\text{size}(\mathcal{N}_\varepsilon) \sim \mathcal{O}\left(\varepsilon^{-\vartheta_\varepsilon}\right) \quad \text{for some } \vartheta_\varepsilon \geq 0.$$

Our DeepONet approximation, with underlying measure μ , is said to *incur a curse of dimensionality* if $\lim_{\varepsilon \rightarrow 0} \vartheta_\varepsilon = +\infty$ and *breaks the curse of dimensionality* if $\lim_{\varepsilon \rightarrow 0} \vartheta_\varepsilon = \bar{\vartheta} < +\infty$.

¹Size of a neural network is understood as the number of non-vanishing parameters of the net.

$$\mathcal{E}_{\text{total}} = \mathcal{E}_{\text{approx}} + \mathcal{E}_{\text{gener}} + \mathcal{E}_{\text{optim}}$$

Definition (Curse of dimensionality)

For a given $\varepsilon > 0$, let $\mathcal{N}_\varepsilon$ be a DeepONet such that $\mathcal{E}_{\text{approx}} < \varepsilon$, and

$$\text{size}(\mathcal{N}_\varepsilon) \sim \mathcal{O}\left(\varepsilon^{-\vartheta_\varepsilon}\right) \quad \text{for some } \vartheta_\varepsilon \geq 0.$$

Our DeepONet approximation, with underlying measure μ , is said to *incur a curse of dimensionality* if $\lim_{\varepsilon \rightarrow 0} \vartheta_\varepsilon = +\infty$ and *breaks the curse of dimensionality* if $\lim_{\varepsilon \rightarrow 0} \vartheta_\varepsilon = \bar{\vartheta} < +\infty$.

Yarotsky proved that the approximation of a general Lipschitz function to accuracy ε requires a ReLU network of size¹ $\varepsilon^{-m(\varepsilon)/2}$, with $m(\varepsilon) \rightarrow \infty$ as $\varepsilon \rightarrow 0$, and hence suffers from the curse of dimensionality.



Yarotsky, D.: *Optimal approximation of continuous functions by very deep relu networks.* **Conference on Learning Theory. PMLR, 639-649, 2018.**

In our setting, m is the number of sensors for the encoding operator $u \mapsto \mathcal{E}(u) = (u(x_1), \dots, u(x_m))$.

¹Size of a neural network is understood as the number of non-vanishing parameters of the net.

However, for some classes of linear and nonlinear operators, the DeepONet approximation may break the curse of dimensionality for **approximation error**.



Lanthaler, S., Mishra, S., Karniadakis, G.E.: *Error estimates for DeepONets: a deep learning framework in infinite dimensions*. **Trans. Math. Appl.** 6(1), 1–144, 2022.

Some examples are:

However, for some classes of linear and nonlinear operators, the DeepONet approximation may break the curse of dimensionality for **approximation error**.



Lanthaler, S., Mishra, S., Karniadakis, G.E.: *Error estimates for DeepONets: a deep learning framework in infinite dimensions*. **Trans. Math. Appl.** 6(1), 1–144, 2022.

Some examples are:

- holomorphic mappings $[-1, 1]^N \rightarrow V$, $y \mapsto \mathcal{G}(y)$ with V an arbitrary Banach space,

However, for some classes of linear and nonlinear operators, the DeepONet approximation may break the curse of dimensionality for **approximation error**.



Lanthaler, S., Mishra, S., Karniadakis, G.E.: *Error estimates for DeepONets: a deep learning framework in infinite dimensions*. **Trans. Math. Appl.** 6(1), 1–144, 2022.

Some examples are:

- holomorphic mappings $[-1, 1]^N \rightarrow V$, $y \mapsto \mathcal{G}(y)$ with V an arbitrary Banach space,
- some PDE operators like
 - 1 parametric elliptic PDEs: $\mathcal{G} : a \mapsto y$, where y solves

$$-\nabla \cdot (a \nabla y) = f$$

However, for some classes of linear and nonlinear operators, the DeepONet approximation may break the curse of dimensionality for **approximation error**.



Lanthaler, S., Mishra, S., Karniadakis, G.E.: *Error estimates for DeepONets: a deep learning framework in infinite dimensions*. **Trans. Math. Appl.** **6**(1), 1–144, 2022.

Some examples are:

- holomorphic mappings $[-1, 1]^N \rightarrow V$, $y \mapsto \mathcal{G}(y)$ with V an arbitrary Banach space,
- some PDE operators like

1 parametric elliptic PDEs: $\mathcal{G} : a \mapsto y$, where y solves

$$-\nabla \cdot (a \nabla y) = f$$

2 parabolic nonlinear PDEs: $\mathcal{G} : y^0 \mapsto y(T)$, where y solves

$$\begin{cases} \frac{\partial y}{\partial t} = \Delta y + f(y) \\ y(0) = y^0 \end{cases}$$

However, for some classes of linear and nonlinear operators, the DeepONet approximation may break the curse of dimensionality for **approximation error**.



Lanthaler, S., Mishra, S., Karniadakis, G.E.: *Error estimates for DeepONets: a deep learning framework in infinite dimensions*. **Trans. Math. Appl.** 6(1), 1–144, 2022.

Some examples are:

- holomorphic mappings $[-1, 1]^N \rightarrow V$, $y \mapsto \mathcal{G}(y)$ with V an arbitrary Banach space,
- some PDE operators like

1 parametric elliptic PDEs: $\mathcal{G} : a \mapsto y$, where y solves

$$-\nabla \cdot (a \nabla y) = f$$

2 parabolic nonlinear PDEs: $\mathcal{G} : y^0 \mapsto y(T)$, where y solves

$$\begin{cases} \frac{\partial y}{\partial t} = \Delta y + f(y) \\ y(0) = y^0 \end{cases}$$

3 scalar conservation laws: $\mathcal{G} : y^0 \mapsto y(T)$, where y solves

$$\begin{cases} \partial_t y + \partial_x (f(y)) = 0 \\ y(0) = y^0 \end{cases}$$

However, for some classes of linear and nonlinear operators, the DeepONet approximation may break the curse of dimensionality for **approximation error**.



Lanthaler, S., Mishra, S., Karniadakis, G.E.: *Error estimates for DeepONets: a deep learning framework in infinite dimensions*. **Trans. Math. Appl.** 6(1), 1–144, 2022.

Some examples are:

- holomorphic mappings $[-1, 1]^N \rightarrow V$, $y \mapsto \mathcal{G}(y)$ with V an arbitrary Banach space,
- some PDE operators like

1 parametric elliptic PDEs: $\mathcal{G} : a \mapsto y$, where y solves

$$-\nabla \cdot (a \nabla y) = f$$

2 parabolic nonlinear PDEs: $\mathcal{G} : y^0 \mapsto y(T)$, where y solves

$$\begin{cases} \frac{\partial y}{\partial t} = \Delta y + f(y) \\ y(0) = y^0 \end{cases}$$

3 scalar conservation laws: $\mathcal{G} : y^0 \mapsto y(T)$, where y solves

$$\begin{cases} \partial_t y + \partial_x (f(y)) = 0 \\ y(0) = y^0 \end{cases}$$

- bounded linear operators $\mathcal{G} : L^2(D) \rightarrow L^2(U)$

Set $\mathbb{T} = [0, \pi]$. Consider the operator

$$\mathcal{G} : L^2(\mathbb{T}) \times H^{-1}(\mathbb{T}) \rightarrow L^2(0, 4\pi), \quad (y^0, y^1) \mapsto \mathcal{G}(y^0, y^1) = u$$

Set $\mathbb{T} = [0, \pi]$. Consider the operator

$$\mathcal{G} : L^2(\mathbb{T}) \times H^{-1}(\mathbb{T}) \rightarrow L^2(0, 4\pi), \quad (y^0, y^1) \mapsto \mathcal{G}(y^0, y^1) = u$$

We take $X = C(\mathbb{T}) \times C(\mathbb{T})$. For the sake of simplicity, we consider the measure $\mu \times \mu$, where μ is the one dimensional measure associated with squared exponential covariance function

$$C(x, x') = e^{-\frac{|x-x'|^2}{2\ell^2}}$$

Set $\mathbb{T} = [0, \pi]$. Consider the operator

$$\mathcal{G} : L^2(\mathbb{T}) \times H^{-1}(\mathbb{T}) \rightarrow L^2(0, 4\pi), \quad (y^0, y^1) \mapsto \mathcal{G}(y^0, y^1) = u$$

We take $X = C(\mathbb{T}) \times C(\mathbb{T})$. For the sake of simplicity, we consider the measure $\mu \times \mu$, where μ is the one dimensional measure associated with squared exponential covariance function

$$C(x, x') = e^{-\frac{|x-x'|^2}{2\ell^2}}$$

The corresponding eigenpairs are given by

$$\phi_k(x) = e^{-ikx}, \quad \lambda_k = \sqrt{2\pi}\ell e^{-(\ell k)^2/2}$$

Set $\mathbb{T} = [0, \pi]$. Consider the operator

$$\mathcal{G} : L^2(\mathbb{T}) \times H^{-1}(\mathbb{T}) \rightarrow L^2(0, 4\pi), \quad (y^0, y^1) \mapsto \mathcal{G}(y^0, y^1) = u$$

We take $X = C(\mathbb{T}) \times C(\mathbb{T})$. For the sake of simplicity, we consider the measure $\mu \times \mu$, where μ is the one dimensional measure associated with squared exponential covariance function

$$C(x, x') = e^{-\frac{|x-x'|^2}{2\ell^2}}$$

The corresponding eigenpairs are given by

$$\phi_k(x) = e^{-ikx}, \quad \lambda_k = \sqrt{2\pi}\ell e^{-(\ell k)^2/2}$$

The key ingredient here is that in the decomposition

$$\begin{array}{ccc} X & \xrightarrow{\mathcal{G}} & Y \\ \downarrow \mathcal{E} & & \uparrow \mathcal{R} \\ \mathbb{R}^{2m} & \xrightarrow{\mathcal{A}} & \mathbb{R}^p \end{array}$$

the error associated with the linear operator $\mathcal{A}$ vanishes since if σ is an ReLU activation function, then

$$\mathcal{A}x + b = \sigma(\mathcal{A}x + b) - \sigma(-(\mathcal{A}x + b)).$$

Theorem (DeepONet approximation of controllability systems)

Consider the control system for $x \in \Omega \subset \mathbb{R}^d$ and $t > 0$:

$$\begin{cases} y' + Ay = 1_\omega u \\ y(0) = y^0 \end{cases}$$

and assume that for some $T > 0$ there exists a unique control $u \in L^2(\omega \times (0, T))$ such that $y(T) = 0$.

Consider the operator $\mathcal{G} : y^0 \rightarrow u$. Then, for any fixed $\varepsilon > 0$ and $\sigma > 0$, there exists a DeepONet $\mathcal{N} = (\beta, \tau)$ achieving an approximation error

$$\mathcal{E}_{\text{approx}} = \|\mathcal{G} - \mathcal{N}\|_{L^2(\mu)} \lesssim \varepsilon$$

and such that

$$\text{size}(\beta) \lesssim \log(\varepsilon^{-1}), \quad \text{depth}(\beta) \lesssim 1, \quad p \sim \log(\varepsilon^{-1})^d, \quad m \sim \log(\varepsilon^{-1})^{2d+\sigma}.$$



García-Cervera, C.J.; Kessler, M.; Pedregal, P.; Periago, F.: *Universal Approximation of Set-Valued Maps and Application to Control*.
Submitted, 2025.

Estimates for generalization error for DeepONet

Boundedness assumption: there exists $\psi : L^2(D) \rightarrow [0, +\infty[$ such that

$$|\mathcal{G}(y)(t)| \leq \psi(y), \quad \sup_{\theta \in [-B, B]^{d_\theta}} |\mathcal{N}_\theta(y)(t)| \leq \psi(y), \quad \forall y \in L^2(D), \forall t \in U,$$

and there exists $C, \kappa > 0$ such that

$$\psi(y) \leq C (1 + \|y\|_{L^2(D)})^\kappa.$$

Lipschitz continuity assumption: there exists $\Phi : L^2(D) \rightarrow [0, +\infty[$ such that

$$|\mathcal{N}_\theta(y)(t) - \mathcal{N}_{\theta'}(y)(t)| \leq \Phi(y) \|\theta - \theta'\|_{\ell^\infty}, \quad \forall y \in L^2(D), \forall t \in U,$$

and there exists $C, \kappa > 0$ such that

$$\Phi(y) \leq C (1 + \|y\|_{L^2(D)})^\kappa.$$

Theorem (Bound for generalization error)

Let $\mathcal{N}_\mathcal{F}$ and $\mathcal{N}_\mathcal{T}$ be as before. If the above two assumptions hold, then

$$\mathbb{E} [\mathcal{L}(\mathcal{N}_\mathcal{T}) - \mathcal{L}(\mathcal{N}_\mathcal{F})] \leq \frac{C}{\sqrt{N}} \left(1 + Cd_\theta \log \left(CB\sqrt{N} \right) \right)^{2\kappa + \frac{1}{2}} \quad (4)$$

where the constant $C = C(\mu, \psi, \Phi)$ is independent of B and d_θ .

Part IV

Numerical Implementation via DeepXDE

$$\begin{cases} -\Delta u(x) = -2d, & x \in D := (0,1)^d \\ u(x) = \sum_{k=1}^d x_k^2, & x \in \partial D \end{cases}$$

This problem has the exact solution

$$u_{\text{exact}}(x) = \sum_{k=1}^d x_k^2$$

$$\begin{cases} -\Delta u(x) = -2d, & x \in D := (0, 1)^d \\ u(x) = \sum_{k=1}^d x_k^2, & x \in \partial D \end{cases}$$

This problem has the exact solution

$$u_{\text{exact}}(x) = \sum_{k=1}^d x_k^2$$

Table: L^2 - relative error $\frac{\|u_{\text{exact}} - u_{\text{PINN}}\|_{L^2(D)}}{\|u_{\text{exact}}\|_{L^2(D)}}$ for different values of the spatial dimension d and number of training points $N = N_{\text{interior}} + N_{\text{boundary}}$

	$N = 100 + 40$	$N = 1000 + 400$	$N = 10000 + 4000$
$d = 2$	9×10^{-4}	1.2×10^{-4}	7.5×10^{-5}
$d = 3$	1.2×10^{-3}	1.4×10^{-4}	1.3×10^{-3}
$d = 5$	2.5×10^{-2}	7.5×10^{-4}	1.9×10^{-4}
$d = 10$	2.2×10^{-1}	4.2×10^{-3}	2.2×10^{-3}
$d = 20$	3.1×10^{-1}	2.6×10^{-2}	2.5×10^{-3}

PINN. Experiment 1: the Laplace-Poisson equation




Figure: Experiment 1. Comparison between PINN (or predicted) solution u_{PINN} (**left**), and exact solution u_{exact} (**right**). Neural network composed of 3 hidden layers and 50 neurons in each layer. No regularization. Number of training points $N_{\text{interior}} = 100$, $N_{\text{boundary}} = 40$.

PINN. Experiment 2: exact control of the wave equation



Compute $u(t)$ such that the solution $y(x, t)$ of the problem

$$\begin{cases} y_{tt} - y_{xx} = 0, & \text{in } (0, 1) \times (0, 2) \\ y(x, 0) = \sin(\pi x), & \text{in } (0, 1) \\ y_t(x, 0) = 0 & \text{in } (0, 1) \\ y(0, t) = 0, & \text{on } (0, 2) \\ y(1, t) = u(t) & \text{on } (0, 2) \end{cases}$$

satisfies

$$y(x, 2) = y_t(x, 2) = 0.$$

Compute $u(t)$ such that the solution $y(x, t)$ of the problem

$$\begin{cases} y_{tt} - y_{xx} = 0, & \text{in } (0, 1) \times (0, 2) \\ y(x, 0) = \sin(\pi x), & \text{in } (0, 1) \\ y_t(x, 0) = 0 & \text{in } (0, 1) \\ y(0, t) = 0, & \text{on } (0, 2) \\ y(1, t) = u(t) & \text{on } (0, 2) \end{cases}$$

satisfies

$$y(x, 2) = y_t(x, 2) = 0.$$

This problem has an exact solution which can be obtained by means of D'Alembert formula. Indeed, by considering the function

$$\tilde{y}^0(x) = \begin{cases} \sin(\pi x) & -1 \leq x \leq 1 \\ 0 & \text{elsewhere,} \end{cases}$$

the explicit exact state is given by

$$y(x, t) = \frac{1}{2} \left(\tilde{y}^0(x - t) + \tilde{y}^0(x + t) \right), \quad 0 \leq x \leq 1, \quad 0 \leq t \leq 2,$$

PINN. Experiment 2: exact control of the wave equation



Compute $u(t)$ such that the solution $y(x, t)$ of the problem

$$\begin{cases} y_{tt} - y_{xx} = 0, & \text{in } (0, 1) \times (0, 2) \\ y(x, 0) = \sin(\pi x), & \text{in } (0, 1) \\ y_t(x, 0) = 0 & \text{in } (0, 1) \\ y(0, t) = 0, & \text{on } (0, 2) \\ y(1, t) = u(t) & \text{on } (0, 2) \end{cases}$$

satisfies

$$y(x, 2) = y_t(x, 2) = 0.$$

This problem has an exact solution which can be obtained by means of D'Alembert formula. Indeed, by considering the function

$$\tilde{y}^0(x) = \begin{cases} \sin(\pi x) & -1 \leq x \leq 1 \\ 0 & \text{elsewhere,} \end{cases}$$

the explicit exact state is given by

$$y(x, t) = \frac{1}{2} \left(\tilde{y}^0(x - t) + \tilde{y}^0(x + t) \right), \quad 0 \leq x \leq 1, \quad 0 \leq t \leq 2,$$

and the exact control is

$$u(t) = \begin{cases} \frac{1}{2} y^0(1 - t) & 0 \leq t \leq 1 \\ -\frac{1}{2} y^0(t - 1) & 1 \leq t \leq 2. \end{cases}$$

PINN. Experiment 2: exact control of the wave equation




Figure: Experiment 1 (linear wave equation). Comparison between exact control $u(t)$ and PINN (or predicted) control $\hat{u}(t; \theta^*)$ (**left**), and error between exact state and PINN state, i.e. $y(x, t) - \hat{y}(x, t; \theta^*)$ (**right**). Neural network composed of 4 hidden layers and 50 neurons in each layer. No regularization. Number of training points $N = 10300$.



$$\begin{aligned} \mathcal{G} : \quad L^2(0,1) \times H^{-1}(0,1) &\rightarrow L^2(0,2) \\ (y^0, y^1) &\mapsto \mathcal{G}(y^0, y^1) := u \end{aligned}$$

where u is the unique control of minimal L^2 -norm of the system

$$\left\{ \begin{array}{ll} y_{tt} - y_{xx} = 0, & \text{in } (0,1) \times (0,2) \\ y(x,0) = y^0(x), & \text{on } (0,1) \\ y_t(x,0) = y^1(x) & \text{on } (0,1) \\ y(0,t) = 0, & \text{on } (0,2) \\ y(1,t) = u(t) & \text{on } (0,2) \\ y(x,2) = y_t(x,2) = 0, & \text{on } (0,1). \end{array} \right.$$



$$\begin{aligned} \mathcal{G} : \quad L^2(0,1) \times H^{-1}(0,1) &\rightarrow L^2(0,2) \\ (y^0, y^1) &\mapsto \mathcal{G}(y^0, y^1) := u \end{aligned}$$

where u is the unique control of minimal L^2 -norm of the system

$$\begin{cases} y_{tt} - y_{xx} = 0, & \text{in } (0,1) \times (0,2) \\ y(x,0) = y^0(x), & \text{on } (0,1) \\ y_t(x,0) = y^1(x) & \text{on } (0,1) \\ y(0,t) = 0, & \text{on } (0,2) \\ y(1,t) = u(t) & \text{on } (0,2) \\ y(x,2) = y_t(x,2) = 0, & \text{on } (0,1). \end{cases}$$

The **matrix of features** corresponds to initial conditions $(y^0(x), y^1(x))$ evaluated at a number of sensor points $x_j \in (0,1)$. It is computed by sampling a Gaussian random field

$$a(x, \omega) = \sum_{i=1}^{\infty} \sqrt{\lambda_i} e_i(x) \xi_i(\omega), \quad (5)$$

where ξ_i are iid standard Gaussian variables, and $\{\lambda_i, e_i(x)\}_{i=1}^{\infty}$ are the eigenvalues and normalized eigenfunctions of the operator

$$\mathcal{C}(\phi)(x) = \int_0^1 C(x, x') \phi(x') dx', \quad C(x, x') = \sigma^2 \exp\left(-\frac{|x - x'|^2}{2\ell^2}\right)$$



The **vector of labels** corresponds to the exact control, which is given by

$$u(t) = \begin{cases} \frac{1}{2}y^0(1-t) + \frac{1}{2}\int_{1-t}^1 y^1(s) ds & 0 \leq t \leq 1 \\ -\frac{1}{2}y^0(t-1) + \frac{1}{2}\int_{t-1}^1 y^1(s) ds & 1 \leq t \leq 2. \end{cases} \quad (6)$$

A data point is a triplet $((y^0, y^1), t, \mathcal{G}(y^0, y^1)(t))$.



The **vector of labels** corresponds to the exact control, which is given by

$$u(t) = \begin{cases} \frac{1}{2}y^0(1-t) + \frac{1}{2}\int_{1-t}^1 y^1(s) ds & 0 \leq t \leq 1 \\ -\frac{1}{2}y^0(t-1) + \frac{1}{2}\int_{t-1}^1 y^1(s) ds & 1 \leq t \leq 2. \end{cases} \quad (6)$$

A data point is a triplet $((y^0, y^1), t, \mathcal{G}(y^0, y^1)(t))$. Moreover, a specific input (y^0, y^1) appears in multiple data points with different values of t . Thus,



The **vector of labels** corresponds to the exact control, which is given by

$$u(t) = \begin{cases} \frac{1}{2}y^0(1-t) + \frac{1}{2}\int_{1-t}^1 y^1(s) ds & 0 \leq t \leq 1 \\ -\frac{1}{2}y^0(t-1) + \frac{1}{2}\int_{t-1}^1 y^1(s) ds & 1 \leq t \leq 2. \end{cases} \quad (6)$$

A data point is a triplet $((y^0, y^1), t, \mathcal{G}(y^0, y^1)(t))$. Moreover, a specific input (y^0, y^1) appears in multiple data points with different values of t . Thus,

$$X_{\text{train}} = \left[\begin{array}{ccc|c} y_{1,1}^{\text{initial}} & \cdots & y_{1,2m}^{\text{initial}} & t_1 \\ \cdots & \cdots & \cdots & \cdots \\ y_{1,1}^{\text{initial}} & \cdots & y_{1,2m}^{\text{initial}} & t_\ell \\ \hline \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots \\ \hline y_{N,1}^{\text{initial}} & \cdots & y_{N,2m}^{\text{initial}} & t_1 \\ \cdots & \cdots & \cdots & \cdots \\ y_{N,1}^{\text{initial}} & \cdots & y_{N,2m}^{\text{initial}} & t_\ell \end{array} \right], \quad y_{\text{train}} = \left[\begin{array}{c} u_1(t_1) \\ \cdots \\ u_1(t_\ell) \\ \hline \cdots \\ \cdots \\ \hline u_N(t_1) \\ \cdots \\ u_N(t_\ell) \end{array} \right]$$

The implementation in DeepXDE is as follows:


```
1 X_train = (X_train[:, :-1], X_train[:, -1:])
2 X_test = (X_test[:, :-1], X_test[:, -1:])
3
4 data = dde.data.Triple(
5     X_train=X_train, y_train=y_train, X_test=X_test, y_test=
6     y_test
7 )
```



Lu, L., Jin, P., Pang, G., Zhang, Z., Karniadakis, G.E.: *Learning nonlinear operators via deeponet based on the universal approximation theorem of operators*. **Nature Machine Intelligence** 3(3), 218-229, 2021.

DeepONet. Experiment 4: fitting the deeponet model

Next, we select the architecture of the DeepONet



DeepONet. Experiment 4: fitting the deeponet model



The implementation in DeepXDE is as follows:

```
1      m = X_train_input_dimension  # inferred from
      training data
2
3      dim_t = 1  # input dimension of the trunk
4
5      p = 10  # output dimension of the branch and trunk
      nets
6
7      net = dde.nn.DeepONet(
8          [m, 40, p],          # dimensions of the branch net
9          [dim_t, 40, p],      # dimensions of the trunk net
10         "relu",              # activation function
11         "Glorot normal",     # initialization of parameters
12     )
13
14     model = dde.Model(data, net)
15
16     model.compile("adam", lr=0.001)
17
18     losshistory, train_state = model.train(iterations=
19         ITERATIONS)
20
21     dde.saveplot(losshistory, train_state, issave=True,
22         isplot=True)
```




Figure: Experiment 4: wave equation. Exact versus predicted solutions for the smooth initial conditions $y^0(x) = y^1(x) = \sin(\pi x)$. $n_{\text{functions}} = 10^4$, $(\ell_0, \ell_1) = (0.25, 0.125)$, $n_{\text{sensors}} = 101$, $p = 40$. Relative error $\approx 1\%$.




Figure: Experiment 4: wave equation. Exact versus predicted solutions for the unsmooth initial conditions $y^0(x) = \begin{cases} 4x, & 0 \leq x \leq 0.5 \\ 0, & 0.5 < x \leq 1 \end{cases}$, $y^1(x) = 0$.

$n_{\text{functions}} = 10^4$, $(\ell_0, \ell_1) = (0.03125, 0.03125)$, $p = 100$. $n_{\text{sensors}} = 11$. Relative error $\approx 4\%$.

- Fundación Séneca (Agencia de Ciencia y Tecnología de la Región de Murcia (Spain)) through the programme for the development of scientific and technical research by competitive groups (21996/PI/22).



f SéNeCa⁽⁺⁾

Agencia de Ciencia y Tecnología
Región de Murcia

- Grants PID2022-141957OA-C22 funded by MCIN/AEI/10.13039/501100011033, by RDF A way of making Europe



Thank you for your attention !